

# (BIG)GRAPHS IN THE CLOUD

## BIG DATA PROCESSING

---

Félix Cuadrado

[felix.cuadrado@qmul.ac.uk](mailto:felix.cuadrado@qmul.ac.uk)

---

Queen Mary University of London  
School of Electronic Engineering and Computer Science

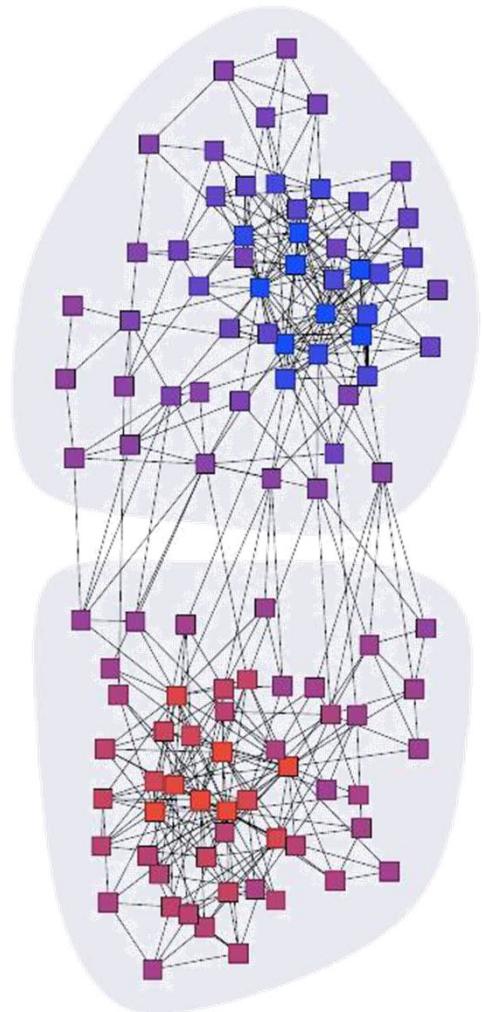
---

# Contents

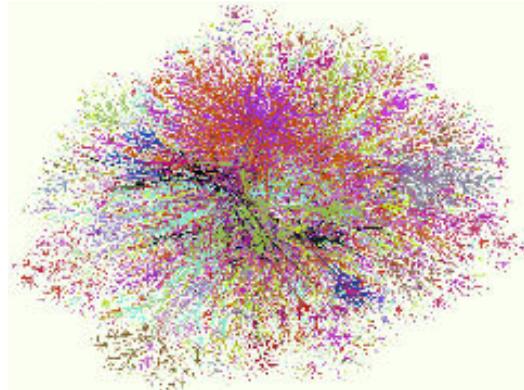
- A world of graphs
- Storing Big Graphs
- Processing Big Graphs

# Graph Definition

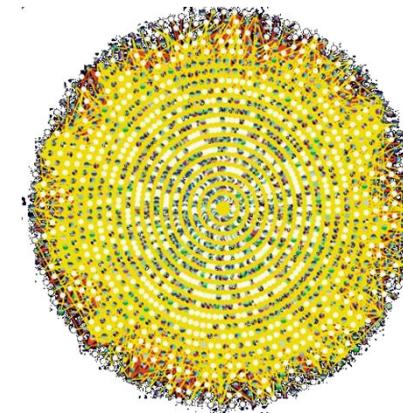
- A graph  $G = (V, E)$ , where
  - $V$  represents the set of vertices (nodes)
  - $E$  represents the set of edges (links)
  - Both vertices and edges may contain additional information
- Different types of graphs:
  - Directed vs. undirected edges
  - Presence or absence of cycles



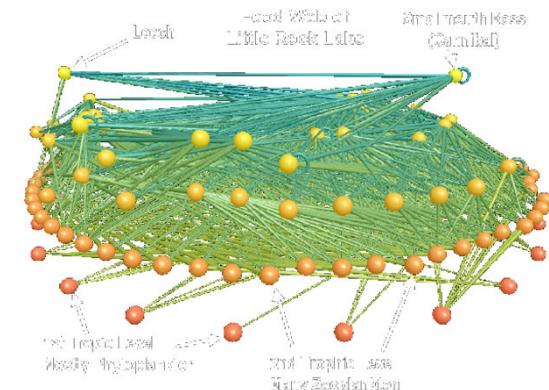
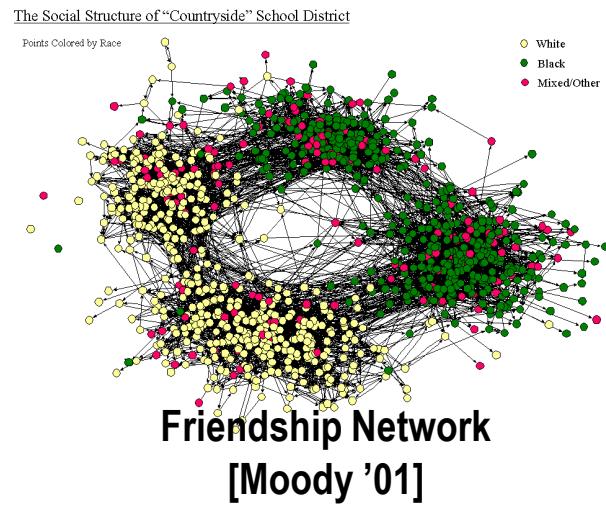
# Graphs are ubiquitous



**Internet Map**  
[lumeta.com]



**Protein Interactions**  
[genomebiology.com]



**Food Web** [Martinez '91]

# Modeling & tracking interactions

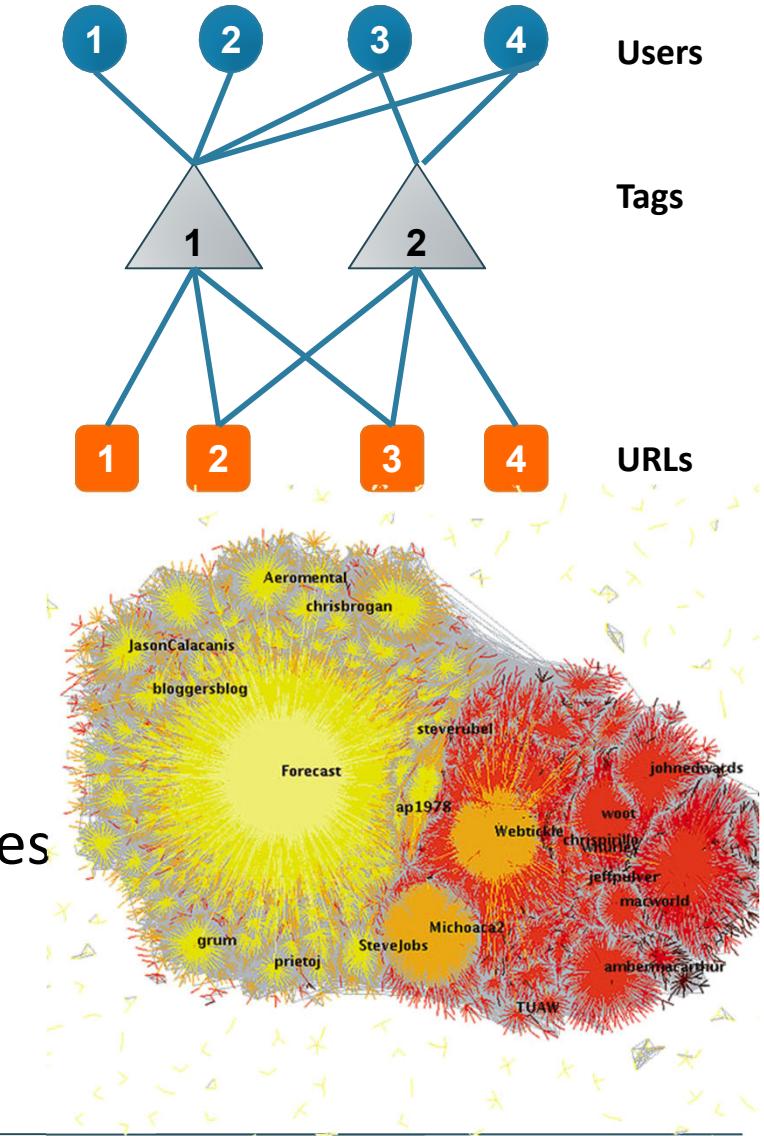


Graphs model  
interactions

Graph size grows  
with input data

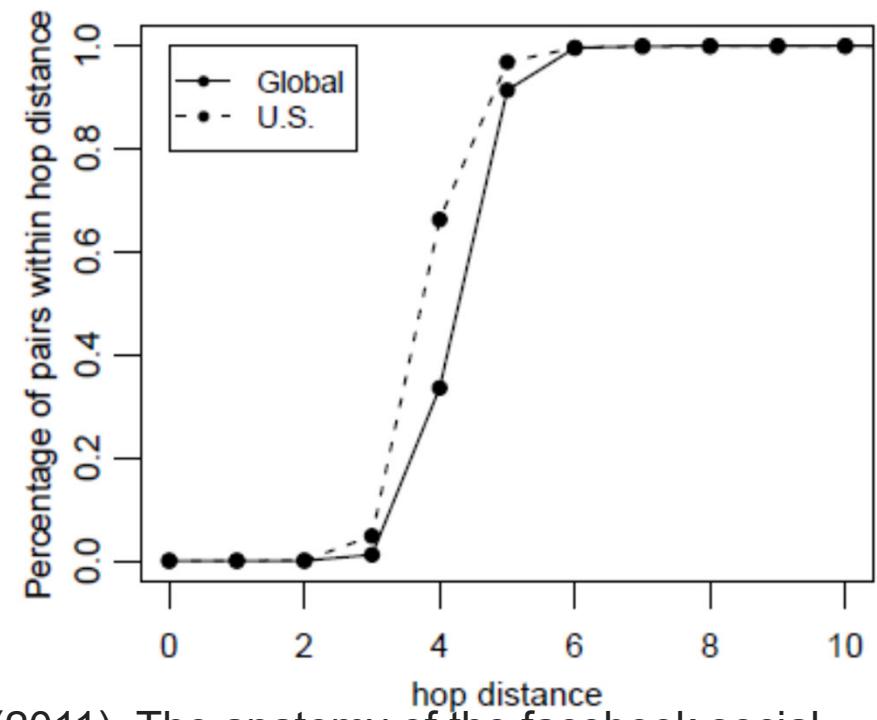
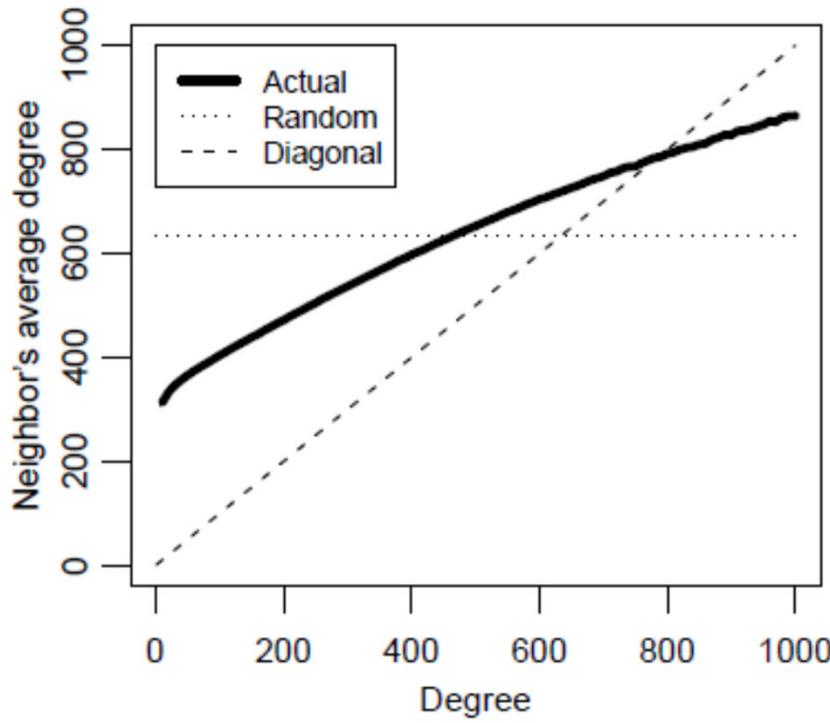
# Social graphs

- Social media defines networks
  - Contacts
  - Messages
  - Tags
- Graph analysis allows to obtain valuable information
  - Identify leaders in a community
    - Measure influence
  - Identify “special” nodes and communities
    - Breaking up terrorist cells
    - Track spread of avian flu



# The Facebook Graph

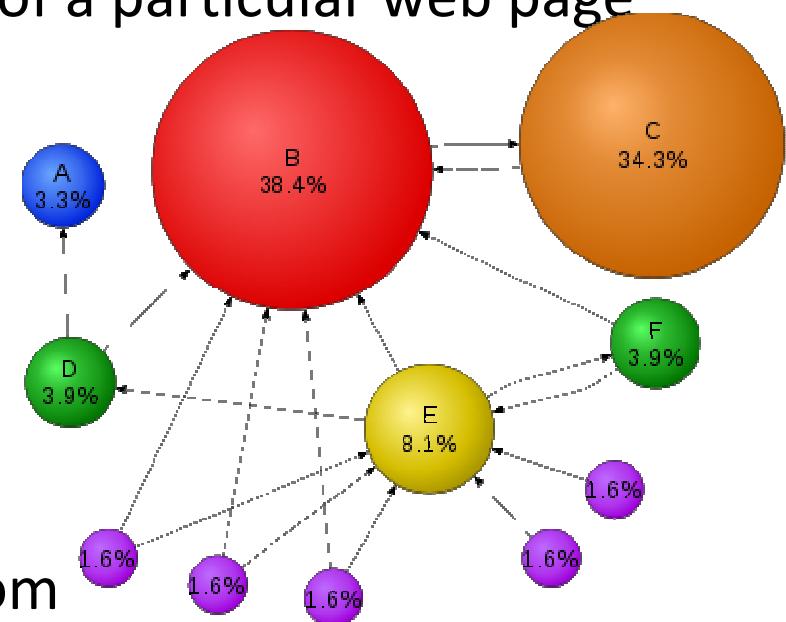
- The power law in the degree popularity
- Replicated in many other human networks



Ugander, J., Karrer, B., Backstrom, L., & Marlow, C. (2011). The anatomy of the facebook social graph. *arXiv preprint arXiv:1111.4503*.

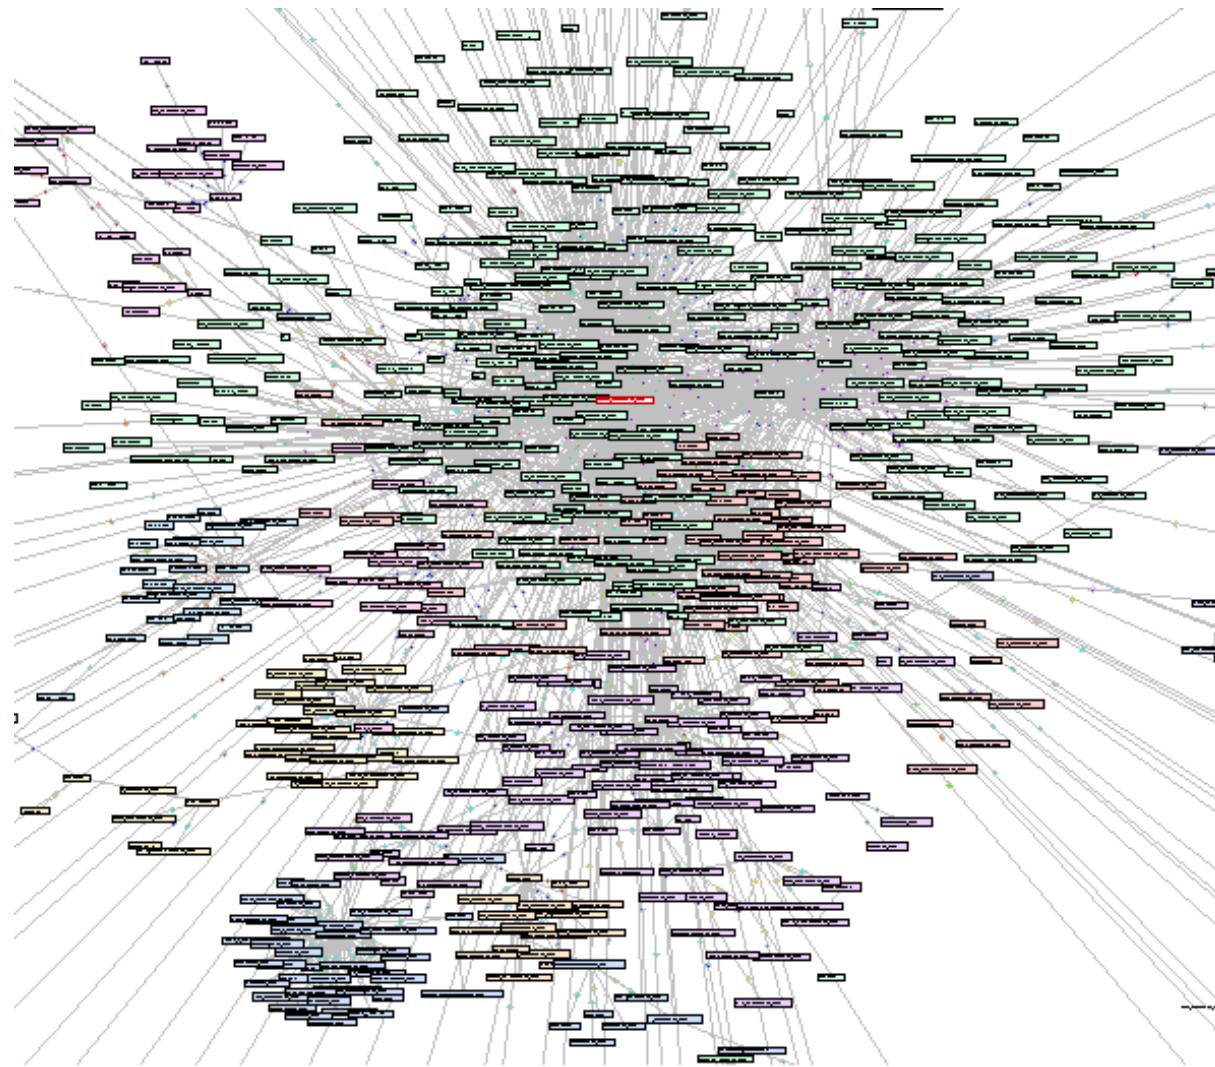
# Google's PageRank

- PageRank is a link analysis algorithm
- The rank value indicates an importance of a particular web page
- A hyperlink to a page counts as a vote of support
- A page that is linked to by many pages with high PageRank receives a high rank itself
- A PageRank of 0.5 means there is a 50% chance that a person clicking on a random link will be directed to the document with the 0.5 PageRank



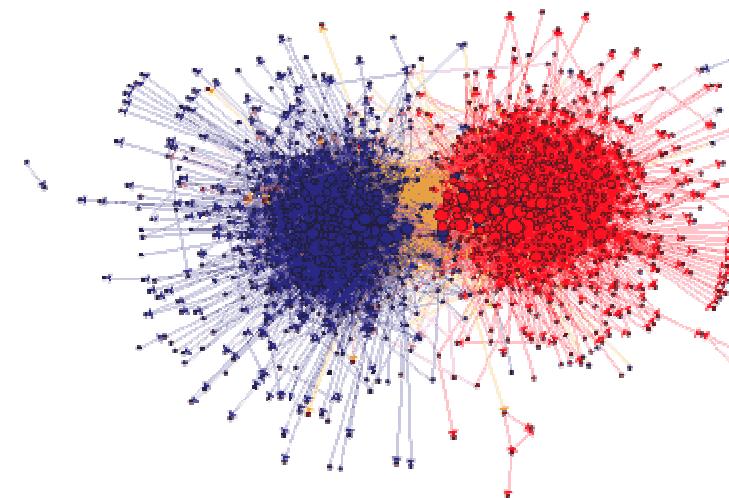
Page, L., Brin, S., Motwani, R., & Winograd, T. (1999). The PageRank citation ranking: bringing order to the web., WWW

# Community detection



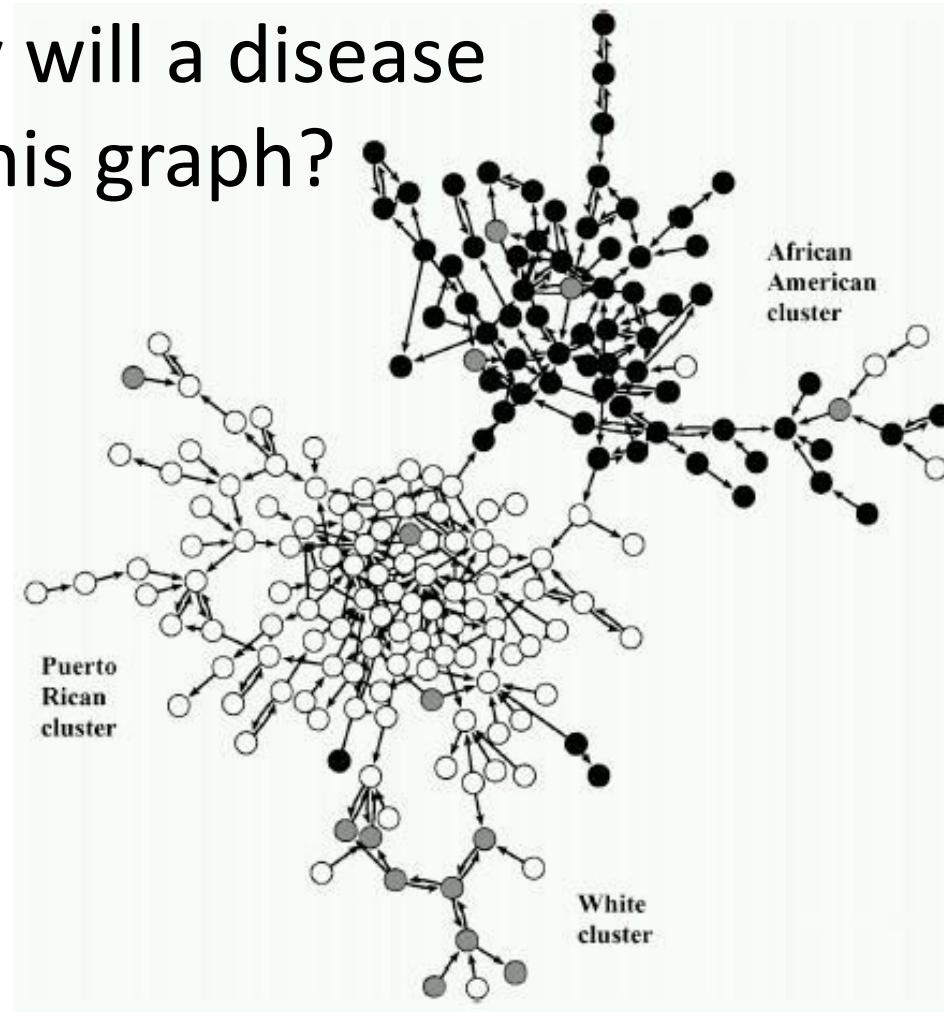
# Bipartite graphs

- Nodes only can have edges to the other part
- “Stable marriage” problem
  - Monster.com, Match.com
- Web advertising, click prediction



# Contagion / epidemic networks

- How quickly will a disease spread on this graph?



# Contents

- A world of graphs
- **Storing Big Graphs**
- Processing Big Graphs

# Graph Management / Storage

- Traditional DBs, NoSQL DBs can store graphs
- But query languages do not support native queries on graph elements (checking for relations)
- We need query languages/abstractions suitable for finding relationship patterns
- Rise of graph DBs
  - Neo4j
  - Titan

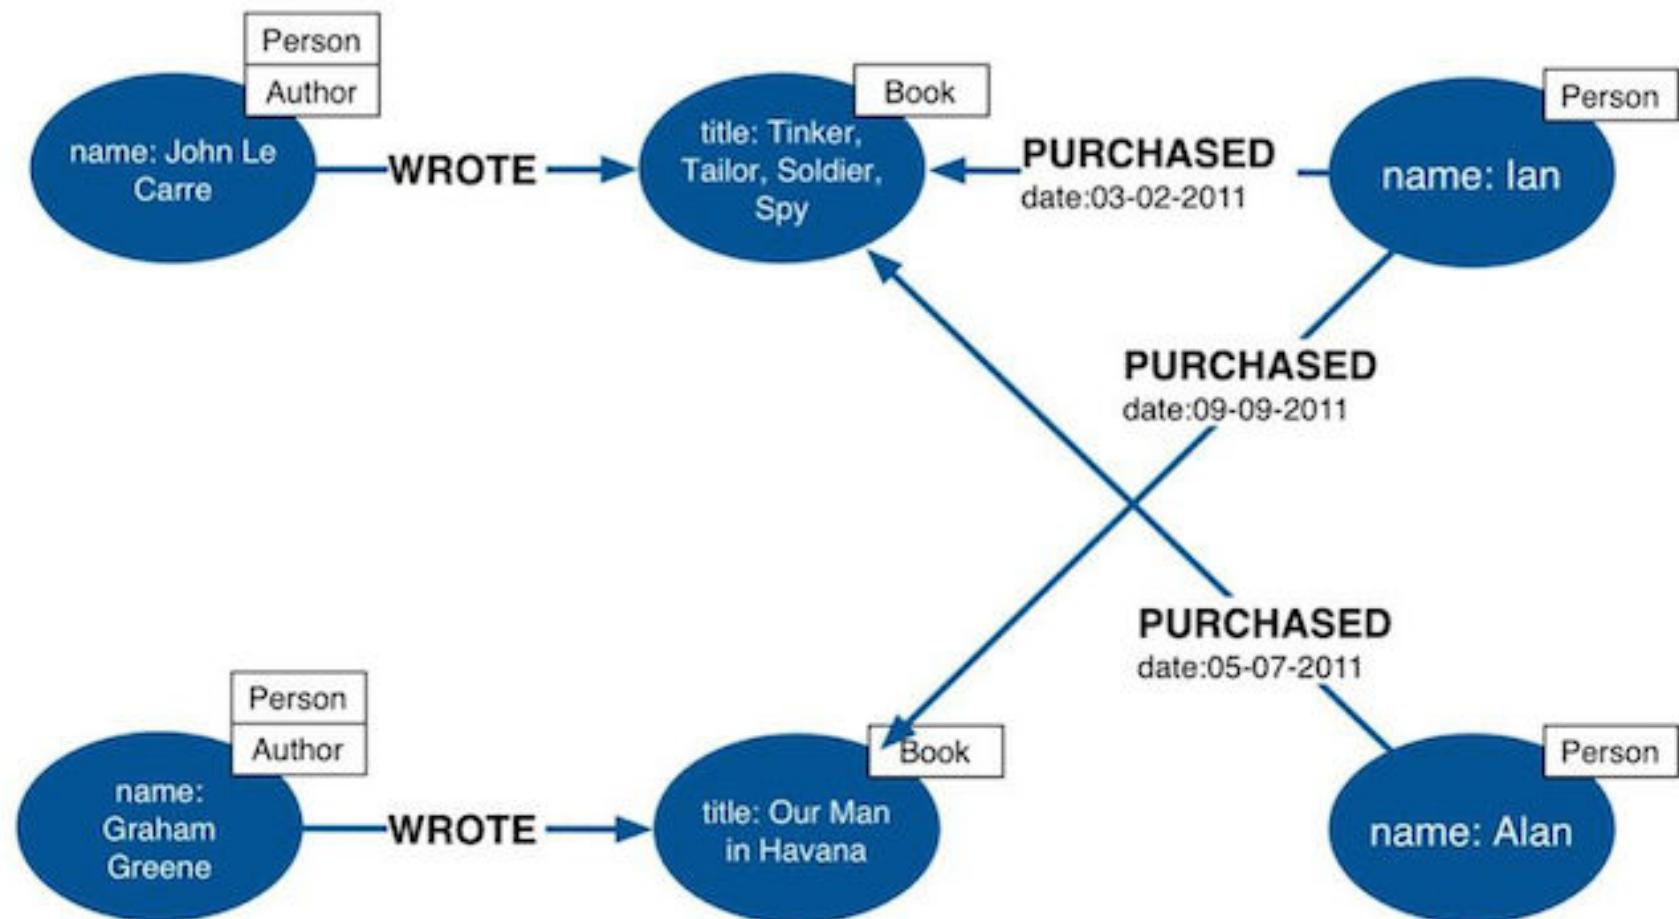
# Graph Databases

- Database that uses graph structures with nodes, edges and properties to store data
- Provides index-free adjacency
  - Every node is a pointer to its adjacent element
  - Fast for following relationships
- Edges hold most of the important information and connect
  - nodes to other nodes
  - nodes to properties

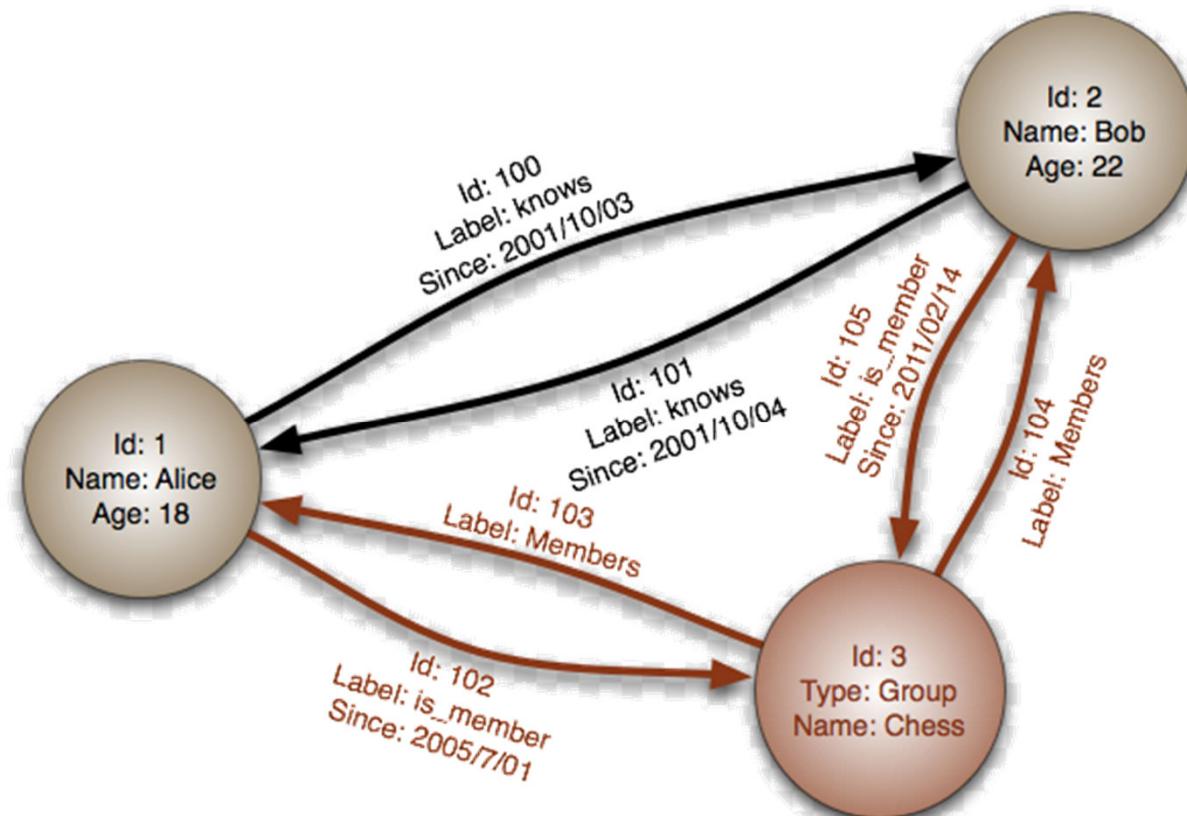
# Neo4j

- Java-based graph database
- ACID – atomic, consistent, isolated and durable for logical units of work
- Property model: powerful schema-less way to model graph-based information
- Good performance for not massive datasets
- Cypher - Graph-specific query language
  - ASCII-art syntax. Define and match patterns

# The property graph model



# Property Graph



# SQL: Modeling and Querying a Graph

Task: listing of all friends-of-a-friend-of-a-friend of mine (who is obviously not a friend of mine)

ACCOUNT HOLDER	FRIEND OF HOLDER
Person001	Person002
Person001	Person006
Person002	Person004
Person004	Person005
Person006	Person007
Person007	Person008

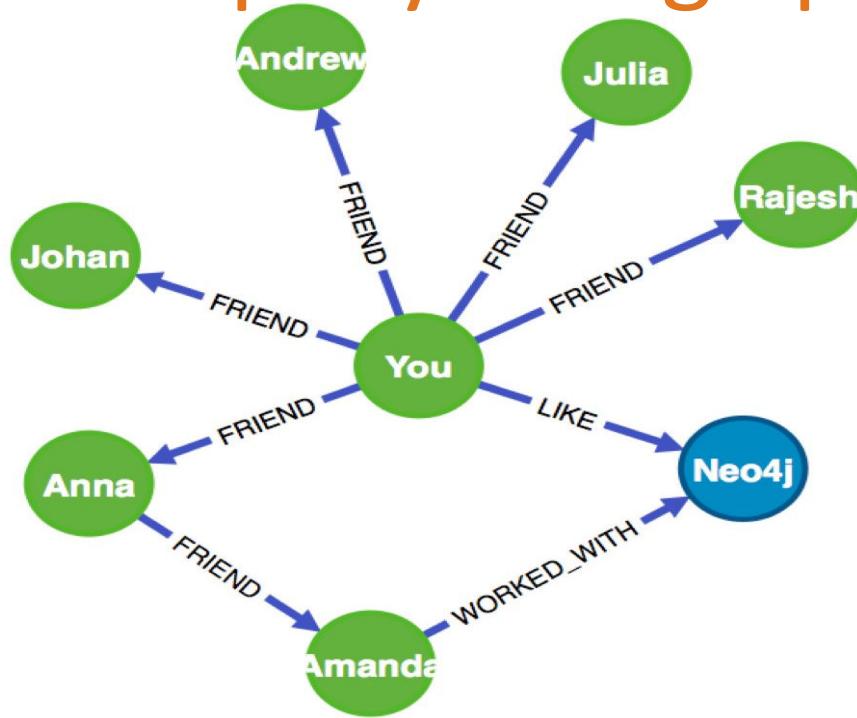
```
SQL: SELECT FRIEND_OF HOLDER FROM TABLE_FRIENDS WHERE FRIEND_OF HOLDER IN
(SELECT FRIEND_OF HOLDER FROM TABLE_FRIENDS WHERE ACCOUNT HOLDER IN
(SELECT FRIEND_OF HOLDER FROM TABLE_FRIENDS WHERE ACCOUNT HOLDER IN
(SELECT FRIEND_OF HOLDER FROM TABLE_FRIENDS WHERE ACCOUNT HOLDER = 'Person001')))
AND FRIEND_OF HOLDER NOT IN (SELECT FRIEND_OF HOLDER FROM TABLE_FRIENDS WHERE
ACCOUNT HOLDER = 'Person001'); FRIEND_OF HOLDER
```

# Cypher query language

- Query Language for Neo4j
  - Becoming standard through OpenCypher initiative
- Match queries, returning all the graph elements who satisfy all the
- Expressive, does not require joining or other SQL

```
MATCH [a:Person {ID: "Person001"}]-[:IS_A_FRIEND_OF*3]->[b:Person]
WHERE NOT (a)-[:IS_A_FRIEND_OF]-(b)
RETURN DISTINCT b.id;
```

# Sample Cypher query on a graph



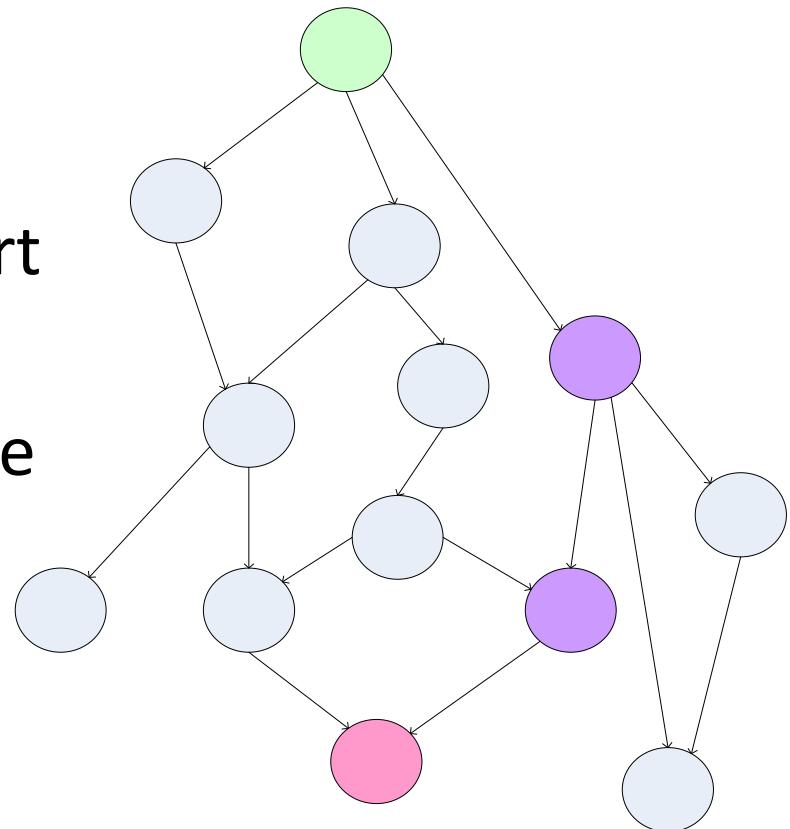
```
MATCH (you {name:"You"})
MATCH (expert)-[:WORKED_WITH]->(db:Database {name:"Neo4j"})
MATCH path = shortestPath( (you)-[:FRIEND*..5]-(expert) )
RETURN db,expert,path
```

# Contents

- A world of graphs
- Storing Big Graphs
- Processing Big Graphs

# Sample graph problem: Finding the Shortest Path

- SSSP: The Single-Source shortest path problem.
- Common graph search application: finding the shortest path from a start node to one or more target nodes
- Commonly done on a single machine with *Dijkstra's Algorithm*



# Dijkstra's algorithm

```

dist[s] ← 0
for all v ∈ V-{s}
  do dist[v] ← ∞
S ← ∅
Q ← V
vertices)
while Q ≠ ∅
do u ← mindistance(Q,dist)

S ← S ∪ {u}
for all v ∈ neighbors[u]
  do if dist[v] > dist[u] + w(u, v)
    then d[v] ← d[u] + w(u, v)
      (if new shortest path found)
      (set new value of shortest path)
      (if desired, add traceback code)

return dist
  
```

**(distance to source vertex is zero)**

**(set all other distances to infinity)**

**(S, the set of visited vertices is initially empty)**

**(Q, the queue initially contains all**

**(while the queue is not empty)**

**(select the element of Q with the min.  
distance)**

**(add u to list of visited vertices)**

**(if new shortest path found)**

**(set new value of shortest path)**

**(if desired, add traceback code)**

# Graph representation

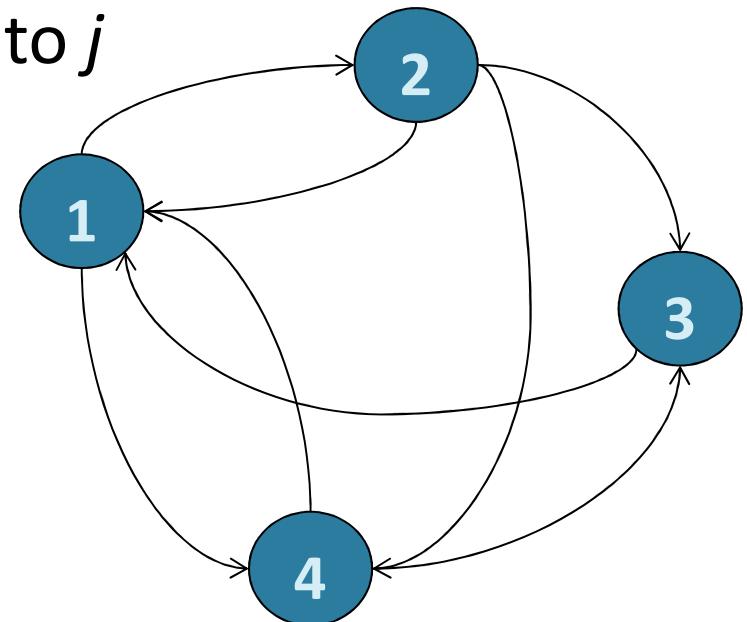
- Algorithms need to keep the information about how graph is represented
  - Structure, and attributes of edges and vertices
  - Both for serialization, and for loading in memory
  - Example: edge list
    - 1 4
    - 2 3
    - 5 2
    - 6 1
    - 1 4

# Graph Representation: Adjacency Matrices

Represent a graph as an  $n \times n$  square matrix  $M$

- $n = |V|$
- $M_{ij} = 1$  means a link from node  $i$  to  $j$

	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	1	0	0	0
4	1	0	1	0



# Adjacency Matrices: Critique

- Advantages:
  - Naturally encapsulates iteration over nodes
  - Rows and columns correspond to inlinks and outlinks
- Disadvantages:
  - Sparse graphs. A graph  $G = (V, E)$  is sparse if  $|E|$  is much smaller than  $|V|^2$ .
  - Lots of zeros for sparse matrices
  - Lots of wasted space

# Graph Traversal in MapReduce

- Approach: Parallel processing of each vertex
  - Each Map/Reduce function has access to limited info
    - One node and its links
- **Iterate** executions of a MapReduce job
  - Map: compute something on each node. Potentially send information to that/other nodes that is aggregated in Reducers.
  - Reducers compute something on each node
  - The output of the reducers in iteration #n becomes the input of the mappers in iteration #n+1

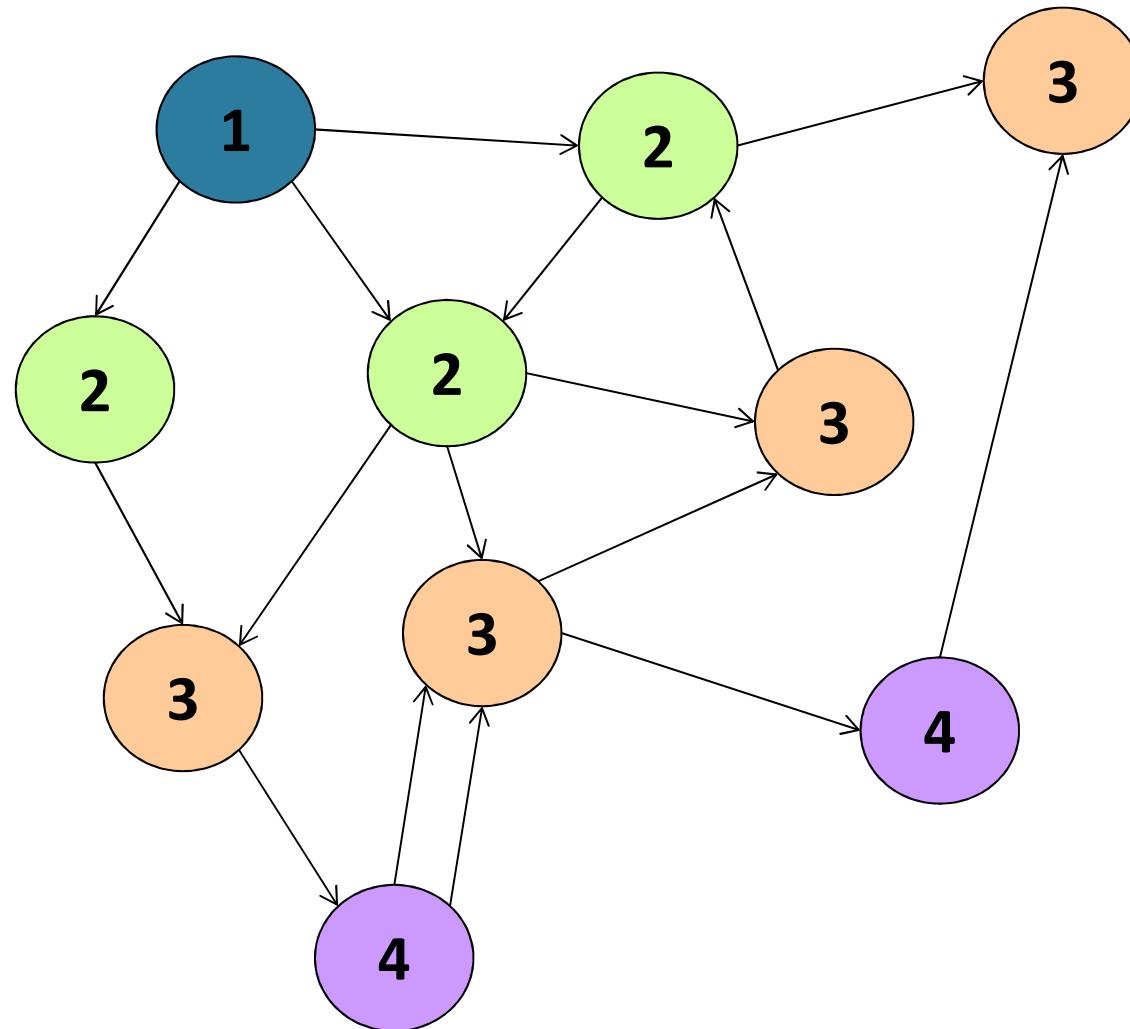
# Equal Weight SSPS

- Problem: from an origin node, find the shortest distance to every other node of the graph, with all links having the same distance
  - Also known as the Kevin Bacon number problem
- Small world phenomenon
  - The world is tied together by a network of personal relationships.
  - There's a popular hypothesis called *Six Degrees of Separation*.

## Finding the Shortest Path: Intuition

- Breadth-First Search (BFS) algorithm
- We can define the solution to this problem inductively:
  - $\text{distanceTo}(\text{startNode}) = 0$
  - For all nodes  $n$  directly reachable from  $\text{startNode}$ ,  
 $\text{distanceTo}(n) = 1$
  - For all nodes  $n$  reachable from some other set of nodes  $S$ ,  
$$\text{distanceTo}(n) = 1 + \min(\text{distanceTo}(m), m \in S)$$

# Visualizing Parallel BFS



# MapReduce graph processing performance

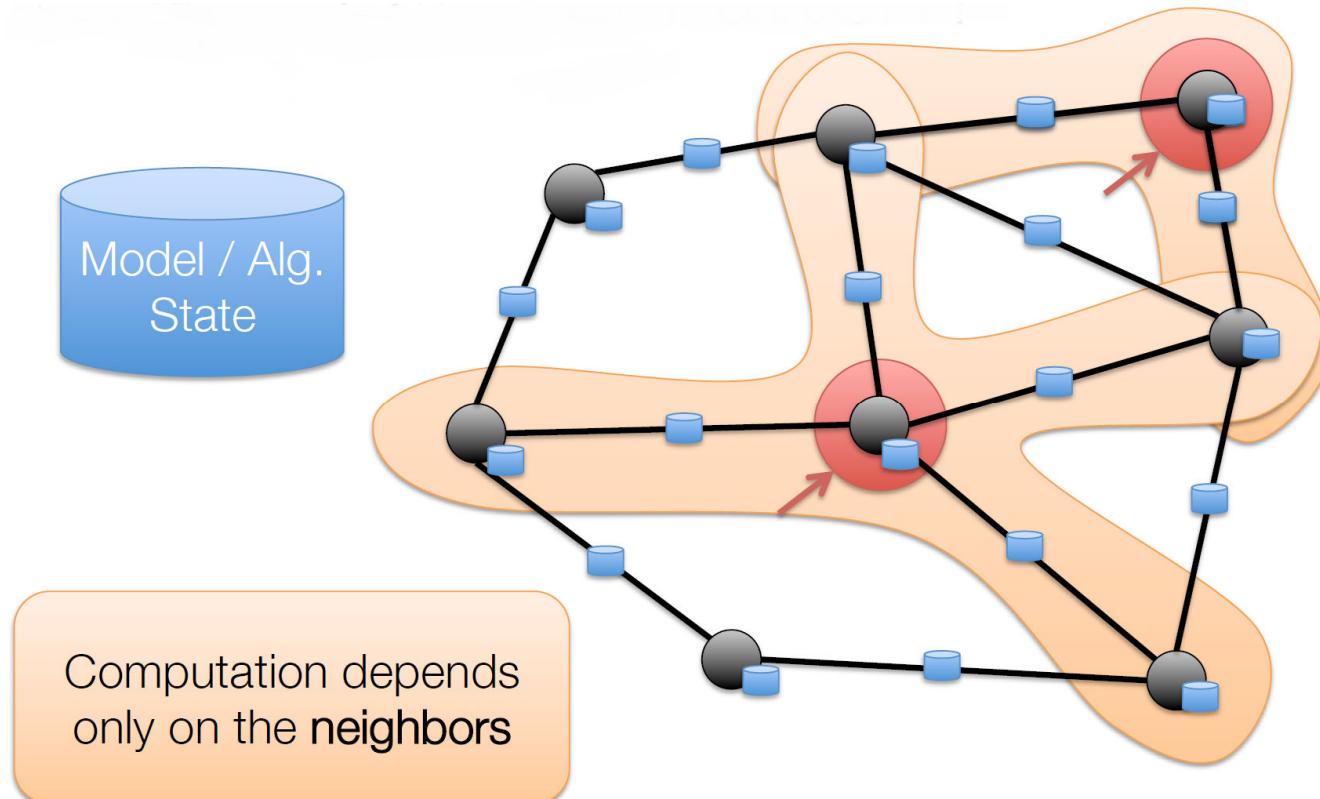
- Iterative algorithms involve HDFS writing in each step
  - Resending the graph structure in each iteration is **VERY inefficient**
- One Map task per node, and sending of messages to other nodes depending on connections results in significant communications cost.
- In-memory systems are a much better fit for this type of computation
  - Graph-specific in-memory systems have developed recently

# Google's Pregel model: Think like a vertex [4]

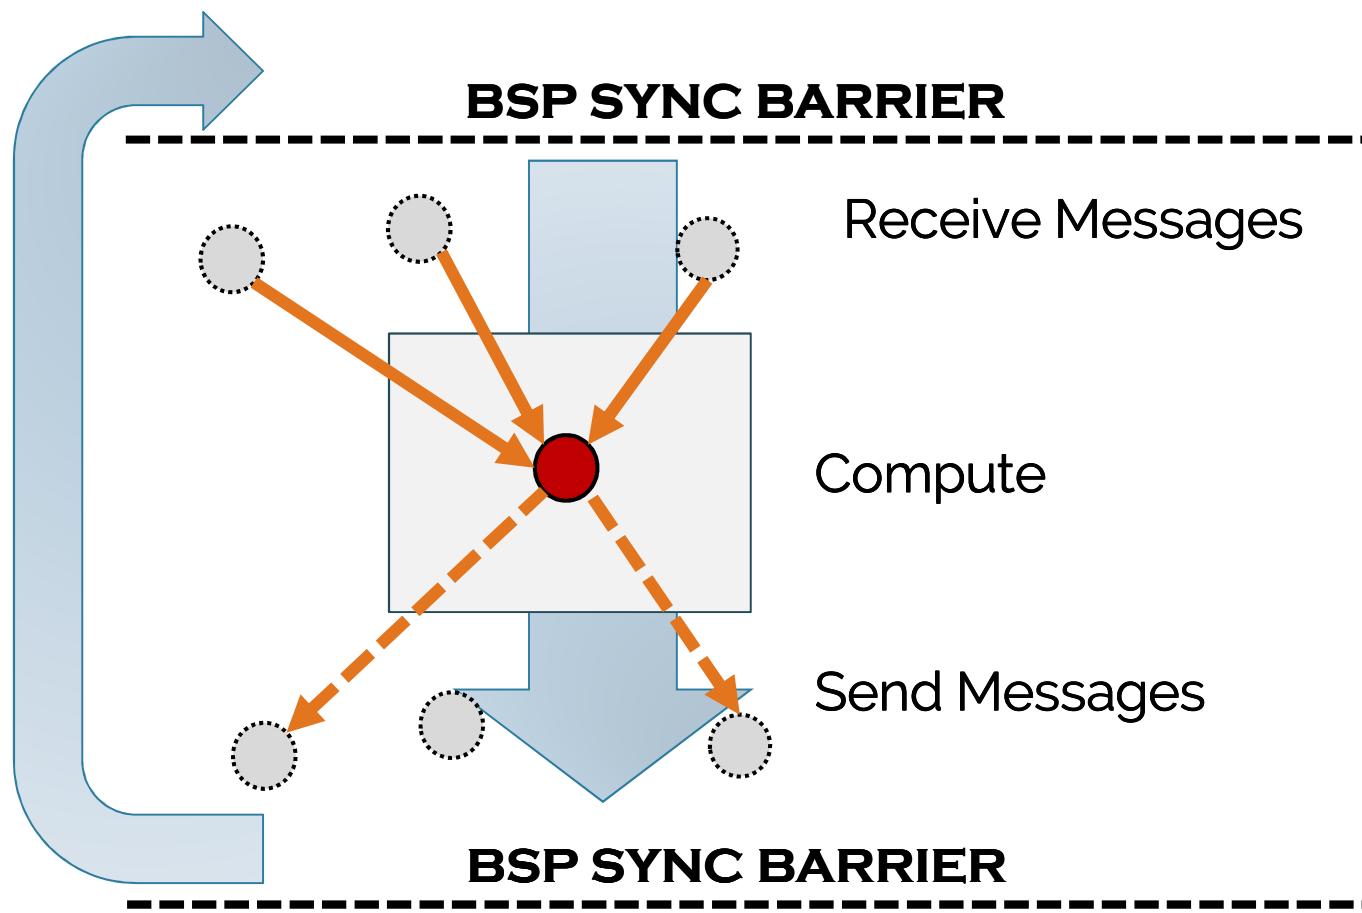
- Computation is iterative
  - Every iteration, a function that is executed at each vertex
    - Vertices can send messages to its neighbours
    - Messages arrive in the next iteration
- Computation is executed in parallel
  - Each vertex is independent from the rest in the same step
  - Messages are the synchronization mechanism

[4] Malewicz, G., Austern, M. H., Bik, A. J., Dehnert, J. C., Horn, I., Leiser, N., & Czajkowski, G. (2010, June). Pregel: a system for large-scale graph processing. In *Proceedings of the ACM SIGMOD*

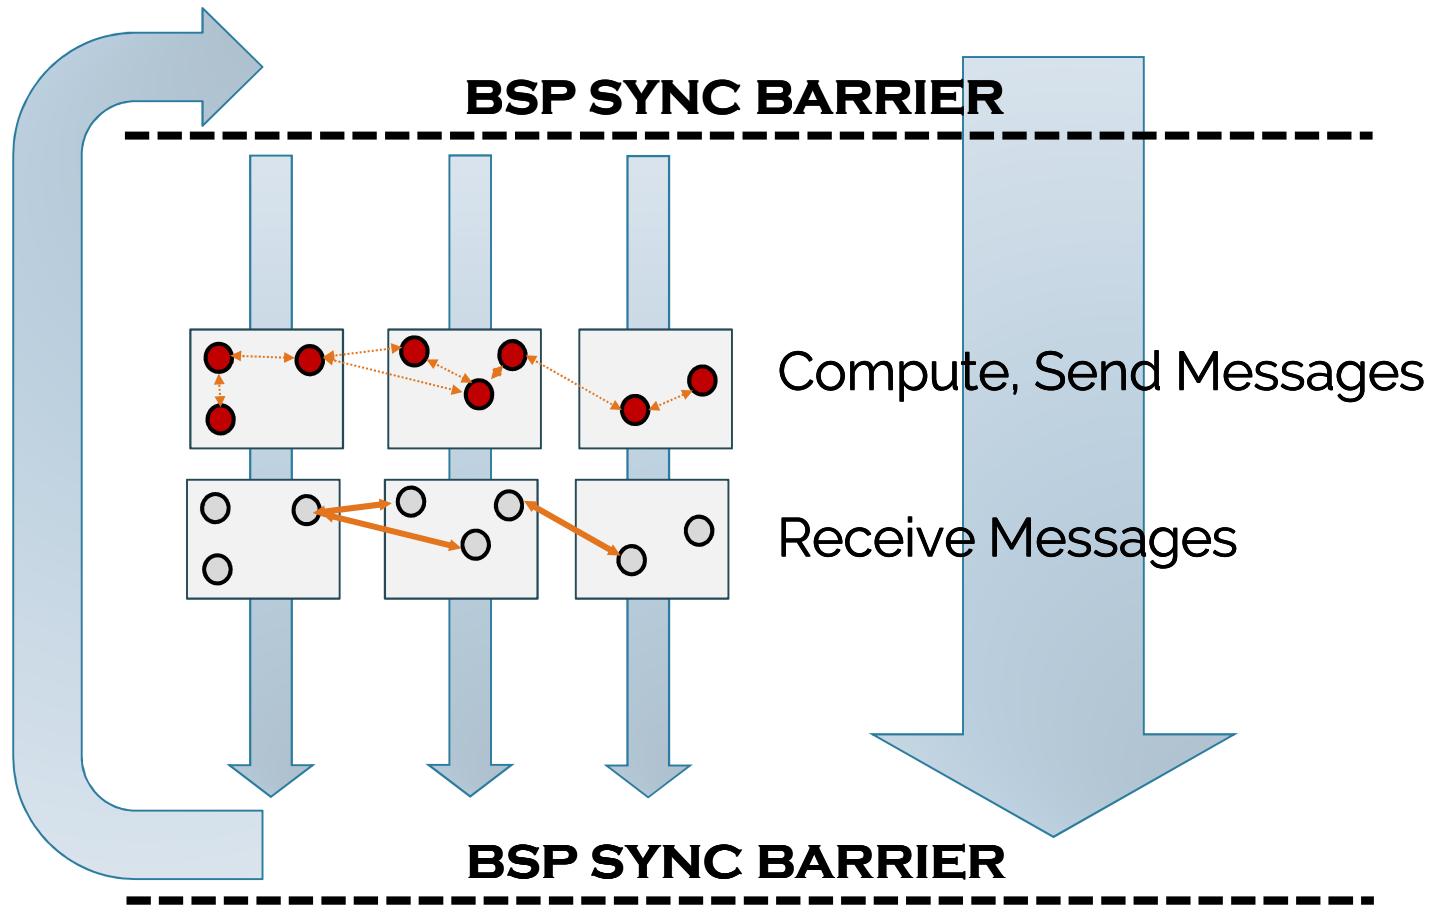
# Pregel: The world of a vertex



# Think like a vertex



# Pregel's node-centric processing model



# Pregel-style graph processing systems

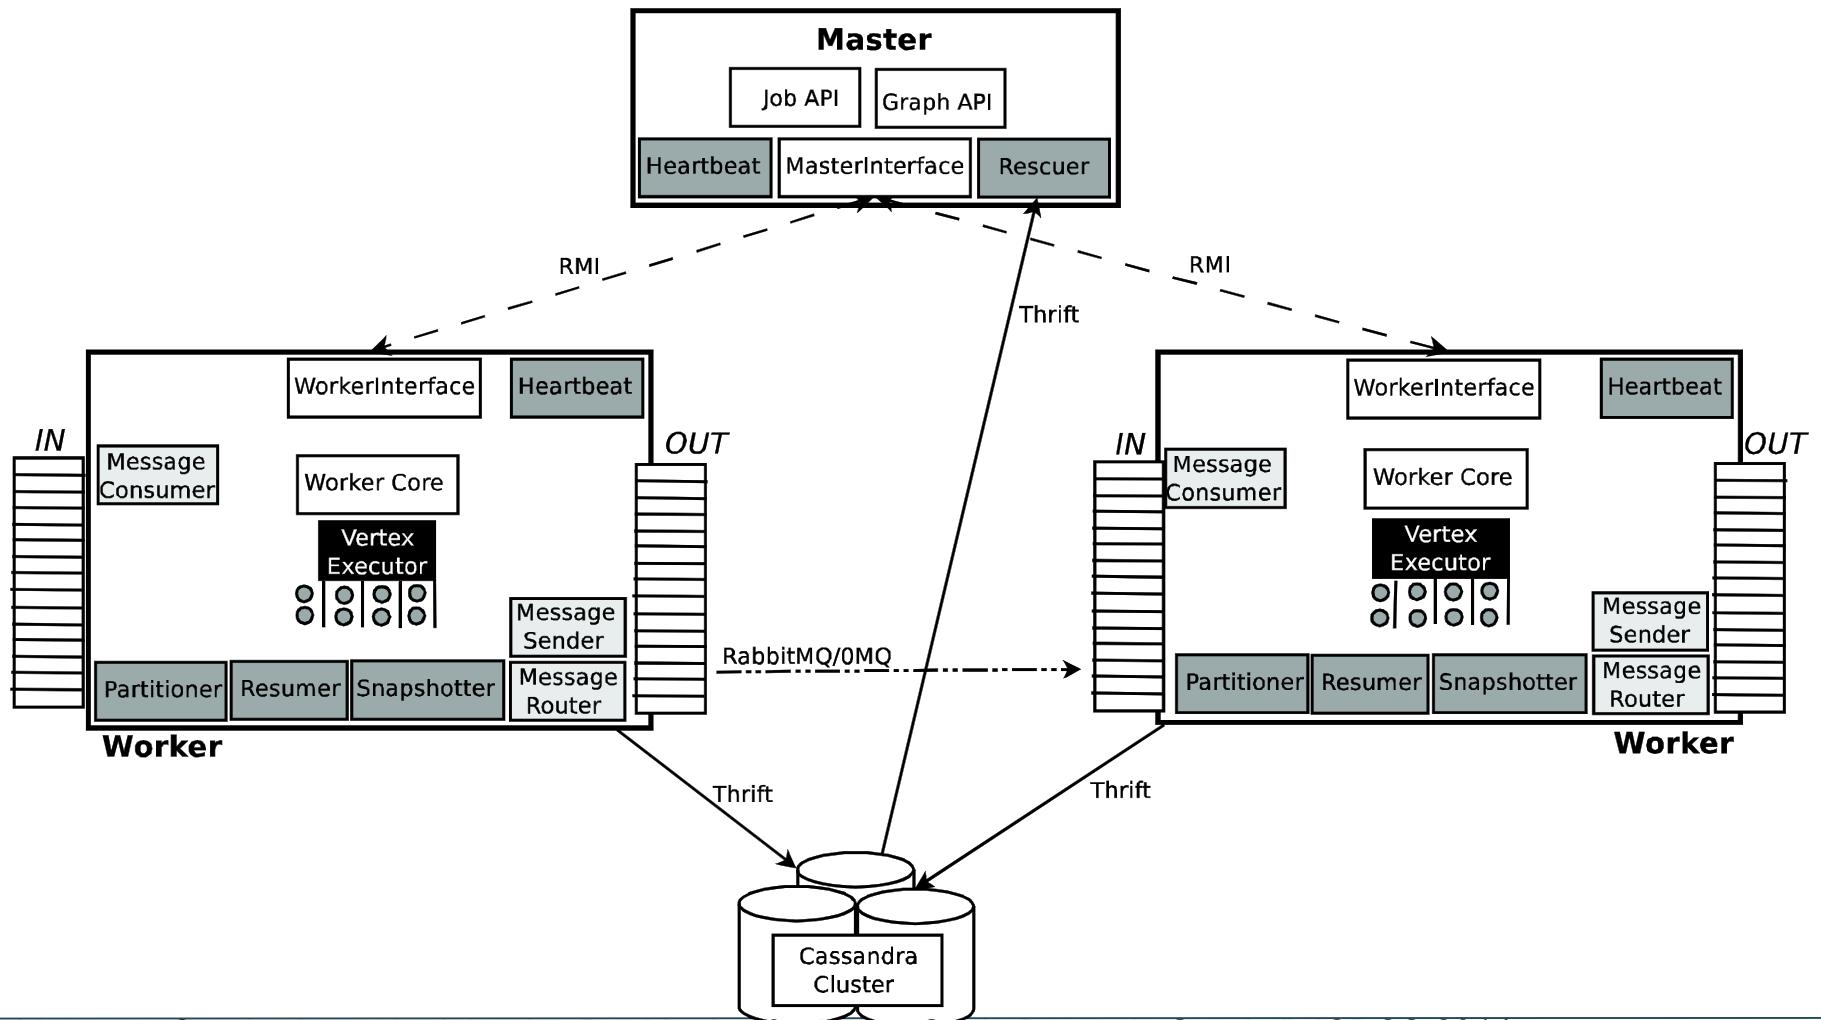
- Pregel
  - Original Google paper
- Apache Giraph
  - Java-based
- Apache Spark GraphX
  - Extension of Spark with
- ...Ongoing research efforts in this space

# Example: PageRank in Apache Giraph

```
public void compute(Vertex <LongWritable, DoubleWritable, FloatWritable> vertex,
                   Iterable messages) throws IOException{
    if (getSuperstep() >= 1) {
        double sum = 0;
        for (DoubleWritable message : messages) {
            sum += message.get();
        }
        DoubleWritable vertexValue =
            new DoubleWritable((0.15f / getTotalNumVertices()) + 0.85f * sum);
        vertex.setValue(vertexValue);
    }

    if (getSuperstep() < MAX_SUPERSTEPS) {
        long edges = vertex.getNumEdges();
        sendMessageToAllEdges(vertex,
            new DoubleWritable(vertex.getValue().get() / edges));
    } else {
        vertex.voteToHalt();
    }
}
```

# GraHPa: Dynamic distributed graph processing

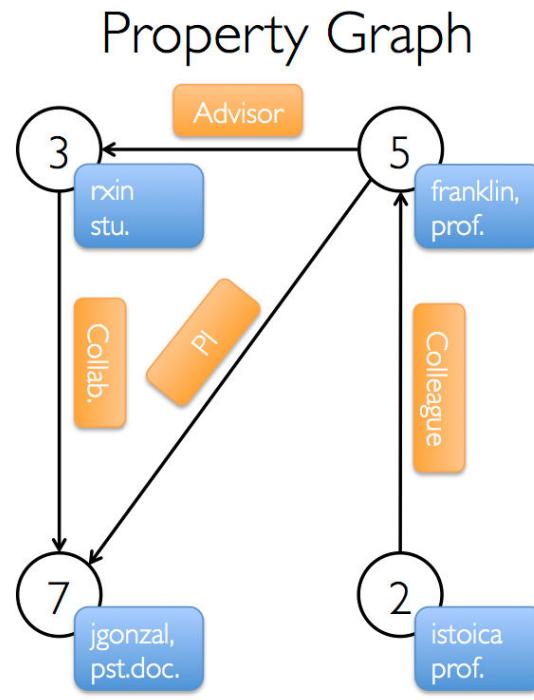


# Spark GraphX

- Spark library for graph processing
- Provides specialized RDDs for representing graph structure, as well as its information (property graphs)
- Provides methods for creating graph, transforming them, implementing multiple common graph metrics and algorithms

# GraphX Property Graphs

```
class Graph[VD, ED]
{ val vertices: VertexRDD[VD]
  val edges: EdgeRDD[ED] }
```



Vertex Table

Id	Property (V)
3	(rxin, student)
7	(jgonzal, postdoc)
5	(franklin, professor)
2	(istoica, professor)

Edge Table

SrcId	DstId	Property (E)
3	7	Collaborator
5	3	Advisor
2	5	Colleague
5	7	PI

```
val userGraph: Graph[(String, String), String]
```

# GraphX RDD

- Holds graph data and provides methods
- VertexRDD[VertexId, VertexData]
  - Vertice IDs have to be longs
    - String to hash conversion
  - Holds vertex properties
- EdgeRDD [EdgeData]
  - Holds source, destination, edge properties
  - Technically a directed multigraph
- Triplets
  - Join of source vertex, destination vertex and edge

# Graph creation in GraphX

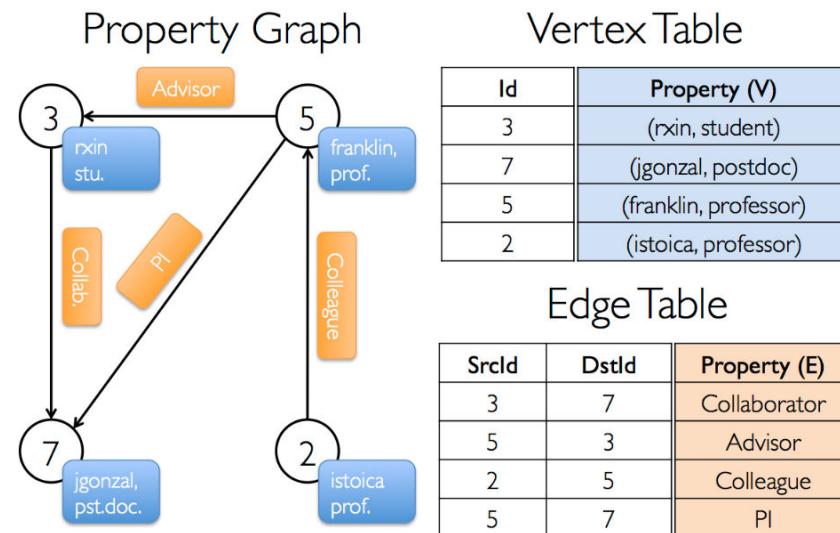
- Graphs can be created from a file, format must be a list of edges in text format. Values will be long numbers
  - `GraphLoader.edgeListFile("hdfs://.....")`
- From RDDs with the right information about nodes and edges.
  - `Graph.fromEdges(<<EdgeRDD>>)`
  - `Graph.fromEdgeTuples(<<RDD[(VertexId, VertexId)] >>)`
- Possible to specify how many partitions, how the graph is partitioned

# GraphX predefined methods

- A graph RDD has multiple convenience methods that provide access to its information and implement relevant operations
- Graph.vertices, graph.edges, graph.triples
  - Access to RDDs with the property information
- Graph.degrees
  - Provides a tuple with (vertexId, degree of each vertex)
- Graph.connectedComponents
  - Obtains each of the connected components of the graph

# Transforming edges and vertices

```
graph.vertices.map(v => v._2._2).collect
//returns (student, postdoc, professor, professor)
graph.edges.filter ( e => e._3.equals("PI").count
//returns 1
```

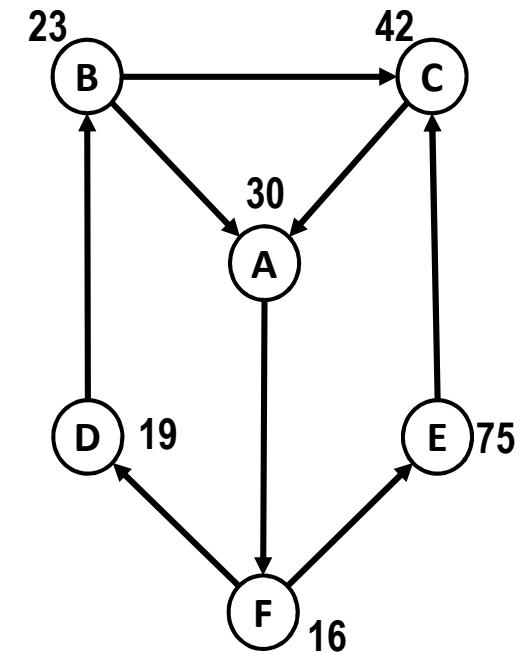


# Graph aggregate computation

- `aggregateMessages` transformation allows to send and process messages through every a message through each edge
  - `sendMsg: EdgeContext[VD, ED, Msg] => Unit`
    - Can send messages to either src or destination, using context
  - `mergeMsg: (Msg, Msg) => Msg`
    - All the receiving messages to a vertex are reduced into one
- Returns a tuple of (`vertexId`, results)

# Age of the oldest follower of each node

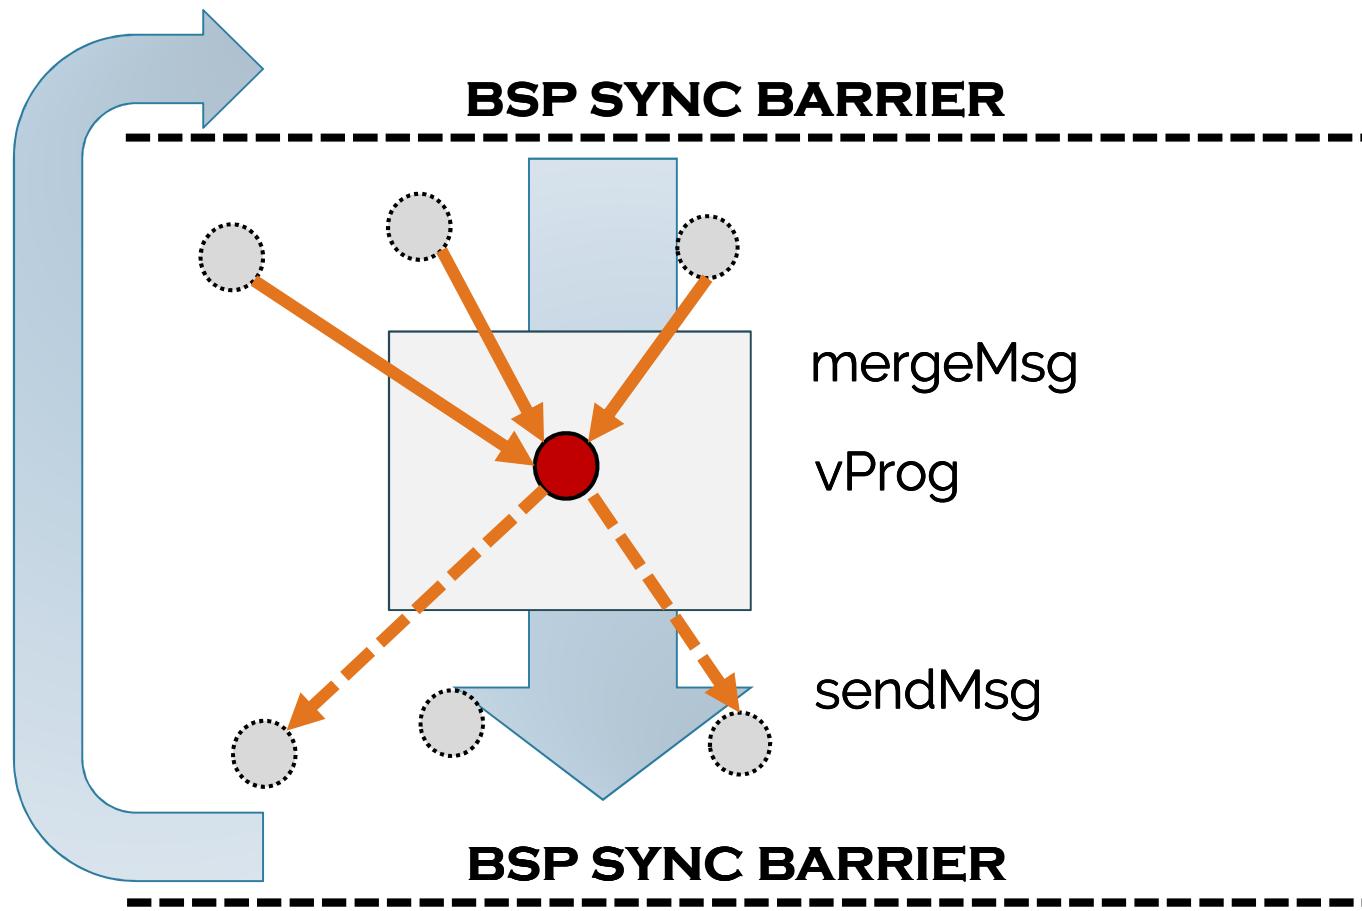
```
val olderFollowers: VertexRDD[(Int, Double)] =  
graph.aggregateMessages[(Int, Double)](  
  // sendMessages  
  edge=> edge.sendToDst(edge.srcAttr) },  
  // mergeMessages  
  (a, b) => math.max(a,b)  
)
```



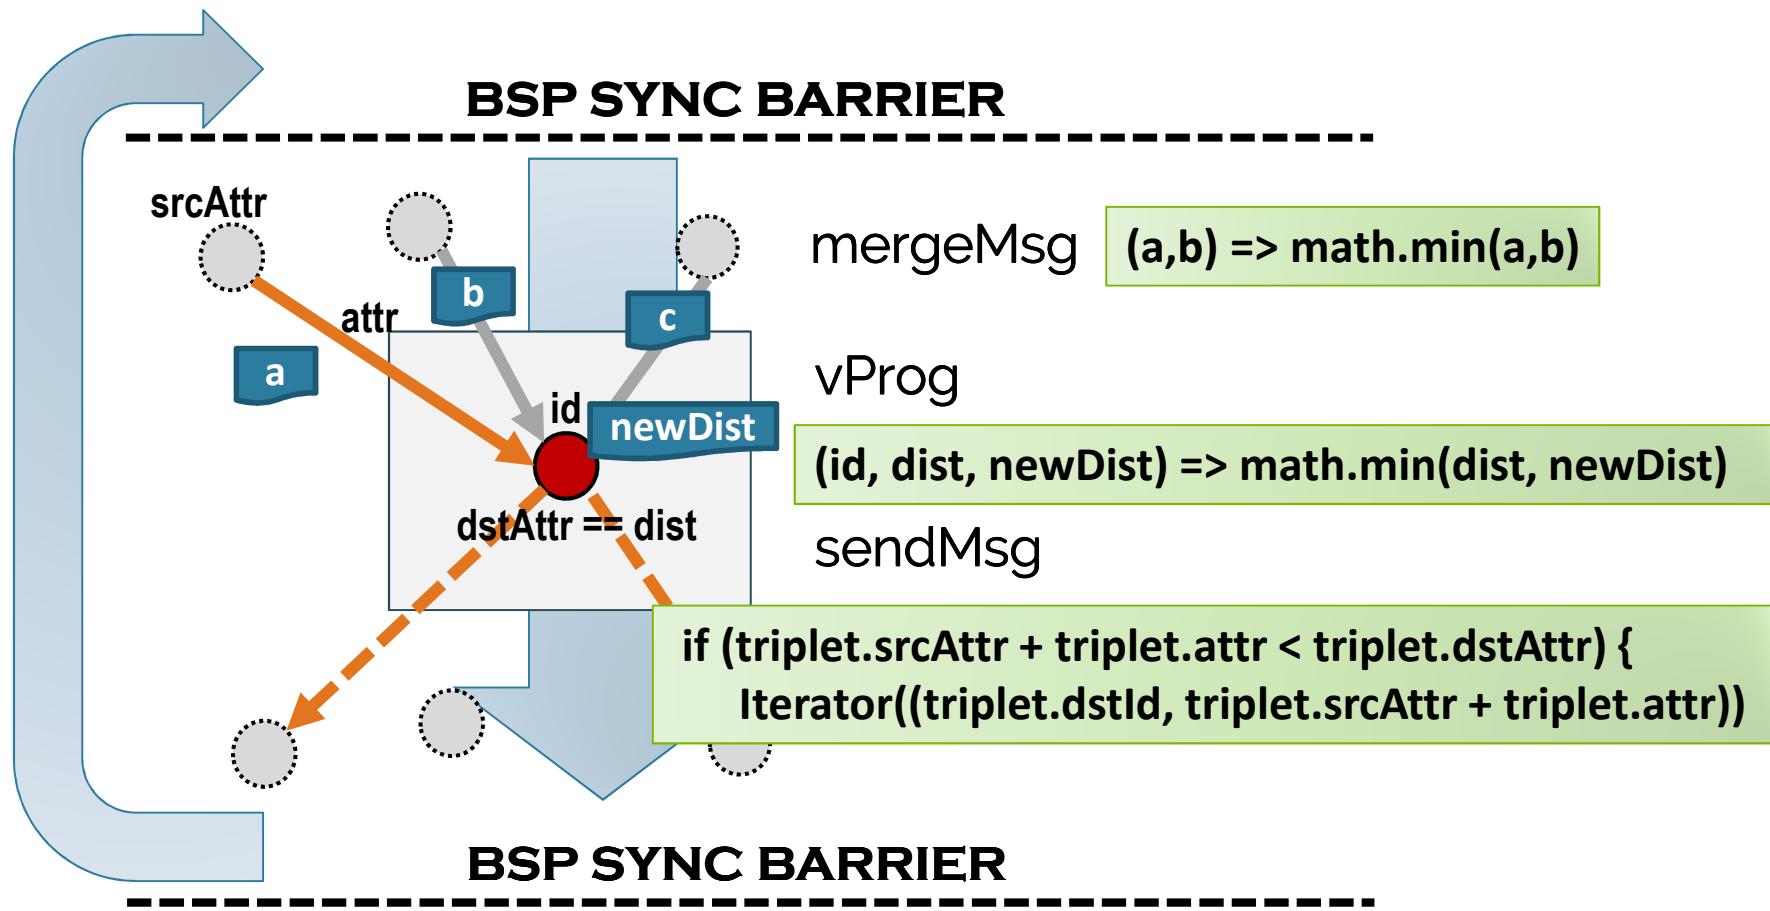
# Pregel programming in GraphX

- GraphX provides a Pregel transformation for iterative graph traversals
- Need to specify number of iterations
- Three functions for implementing Pregel's model
  - `vprog`:  $(\text{VertexID}, \text{VD}, \text{A}) \Rightarrow \text{VD}$ 
    - Updates vertex properties (`VD`)
  - `sendMsg`:  $\text{EdgeTriplet}[\text{VD}, \text{ED}] \Rightarrow \text{Iterator}[(\text{VertexID}, \text{A})]$ 
    - Sends messages to vertex neighbours
  - `mergeMsg`:  $(\text{A}, \text{A}) \Rightarrow \text{A}$ 
    - Combines all incoming messages into a single one
- Returns new graph with updated vertex values

# GraphX Pregel model



# GraphX Pregel model (2)



# Breadth First Search in Pregel

```
val initialGraph = graph.mapVertices((id, _) => if (id == sourcId) 0.0 else Double.PositiveInfinity)
val sssp = initialGraph.pregel(Double.PositiveInfinity)(
  (id, dist, newDist) => math.min(dist, newDist), // Vertex Program
  triplet => { // Send Message
    if (triplet.srcAttr + triplet.attr < triplet.dstAttr) {
      Iterator((triplet.dstId, triplet.srcAttr + triplet.attr))
    } else {
      Iterator.empty
    }
  },
  (a,b) => math.min(a,b) // Merge Message
)
```