

# STREAM PROCESSING

## BIG DATA PROCESSING

---

**Félix Cuadrado**

`felix.cuadrado@qmul.ac.uk`

Queen Mary University of London

School of Electronic Engineering and Computer Science

---

# Contents

- **Information Streams**
- Stream Processing - Storm
- Micro batch processing - DStream

# Information streams


- Data is continuously generated from multiple sources
  - Messages from a social platform (e.g. Twitter)
  - Network traffic going over a switch
  - Readings from distributed sensors
  - Interactions of users with a web application
- For faster analytics, we might need to process the information the moment it is generated
  - Process the information streams

# Streams – A Brave New World

- Batch processing: data stored in *finite, persistent data sets*
- Data Streams: distributed, continuous, unbounded, rapid, time varying, noisy, ...
- Data-Stream Management: variety of modern applications
  - Network monitoring and traffic engineering
  - Sensor networks
  - Telecom call-detail records
  - Network security
  - Financial applications
  - Manufacturing processes
  - Web logs and clickstreams
  - Other massive data sets...


## Example: IP Network Data

- Networks are sources of massive data: the metadata per hour per IP router is gigabytes
- Fundamental problem of data stream analysis: *Too much information to store or transmit*
- So process data as it arrives – *One pass, small space: the data stream approach*
- *Approximate answers* to many questions are OK, if there are guarantees of result quality



# IP Network Monitoring Application

- 24x7 IP packet/flow data-streams at network elements
- Truly massive streams arriving at rapid rates
  - DEC-IX: 5Tb/sec peak traffic
- Often shipped off-site to data warehouse for off-line analysis

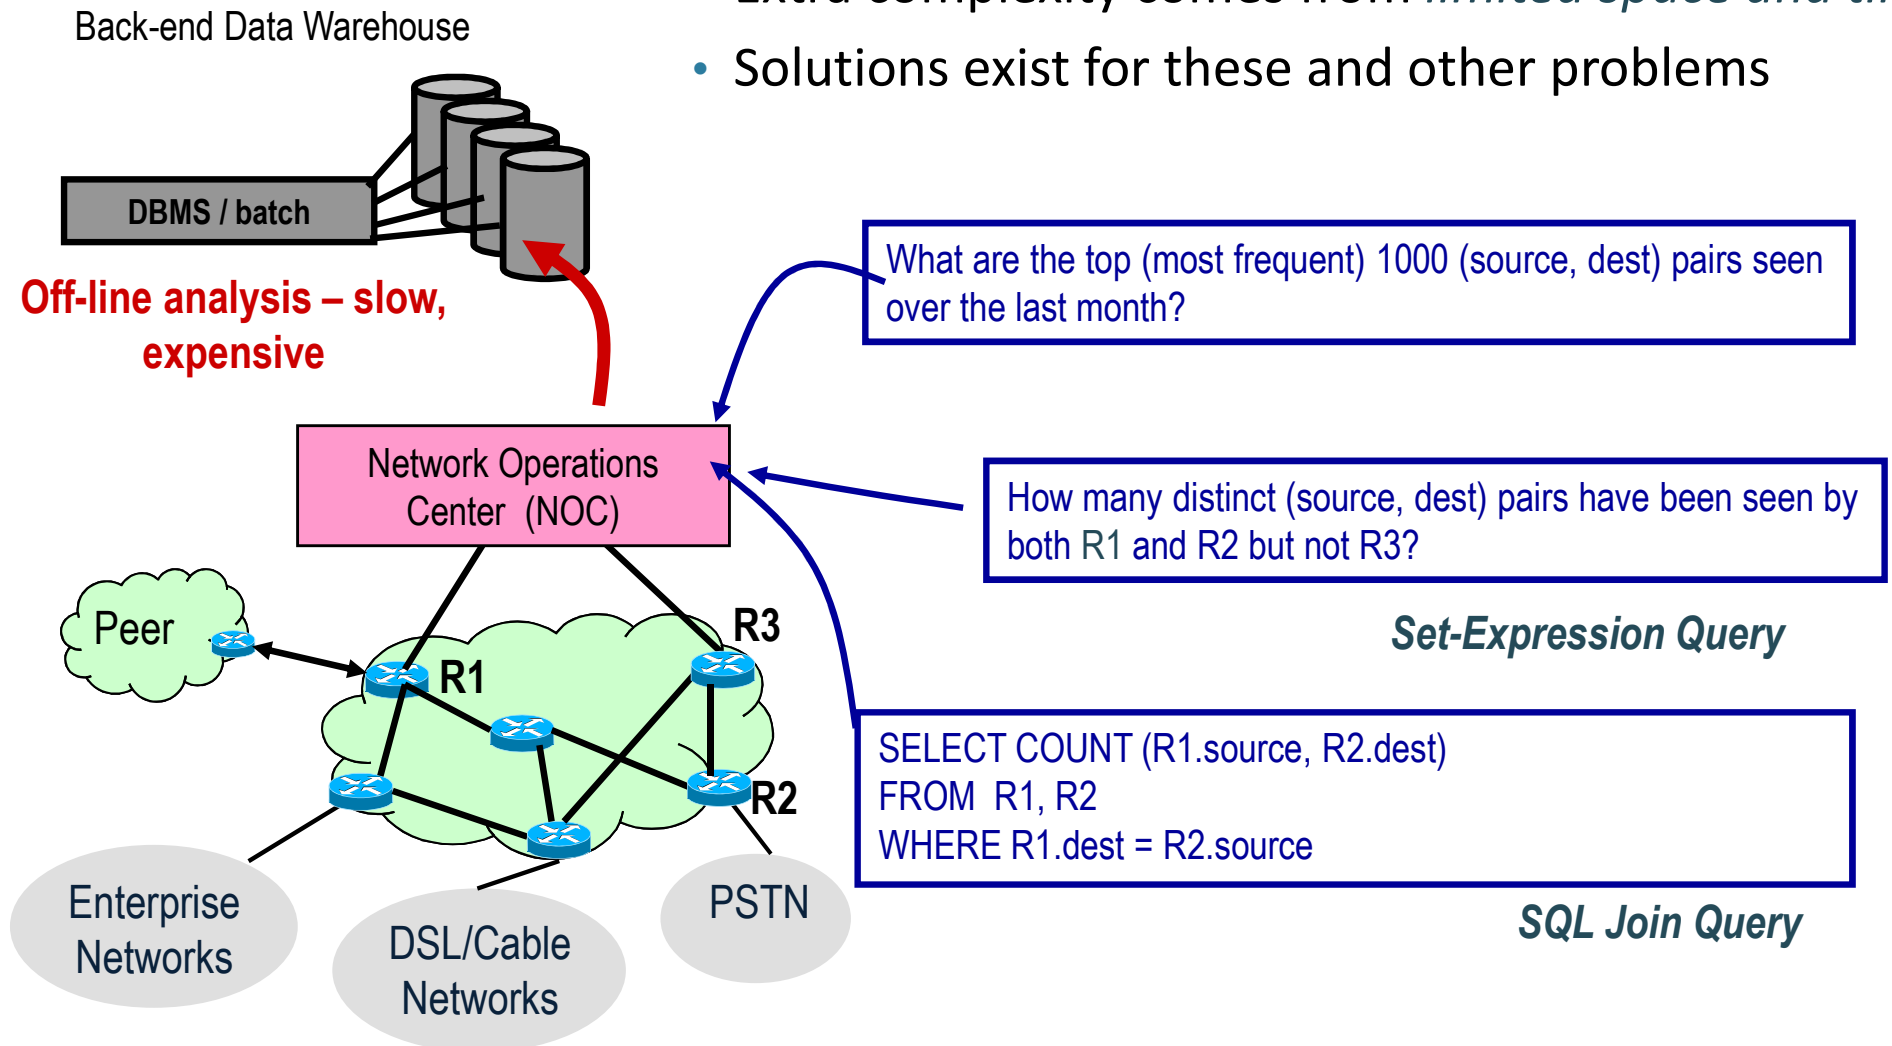


Source	Destination	Duration	Bytes	Protocol
10.1.0.2	16.2.3.7	12	20K	http
18.6.7.1	12.4.0.3	16	24K	http
13.9.4.3	11.6.8.2	15	20K	http
15.2.2.9	17.1.2.1	19	40K	http
12.4.3.8	14.8.7.4	26	58K	http
10.5.1.3	13.0.0.1	27	100K	ftp
11.1.0.6	10.3.4.5	32	300K	ftp
19.7.1.2	16.5.5.8	18	80K	ftp

**Example NetFlow IP Session Data**


# Network Monitoring Queries

- Extra complexity comes from *limited space and time*
- Solutions exist for these and other problems



# Real-Time Data-Stream Analysis

- Must process network streams in *real-time* and *one pass*
- Critical NM tasks: fraud, DoS attacks, SLA violations
  - Real-time traffic engineering to improve utilization
- *Tradeoff result accuracy vs. space/time/communication*
  - Fast responses, small space/time
  - Minimize use of communication resources



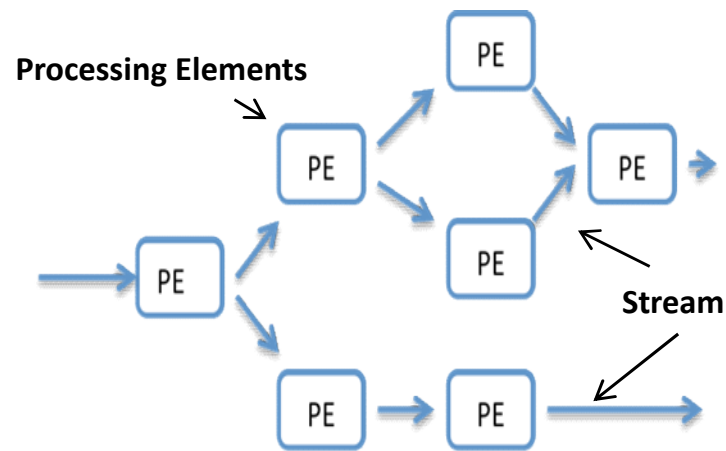


# Contents


- Information Streams
- **Stream Processing - Storm**
- Micro batch processing - DStream

# Stream Processing

- Stream – Sequence of unbounded tuples



Macro view



Microscopic View


# Apache Storm

- Developed by BackType which was acquired by Twitter. Now donated to Apache foundation
- Storm provides realtime computation of data streams
  - Scalable (distribution of blocks, horizontal replication)
  - Guarantees no data loss
  - Extremely robust and fault-tolerant
  - Programming language agnostic

# Storm Concepts

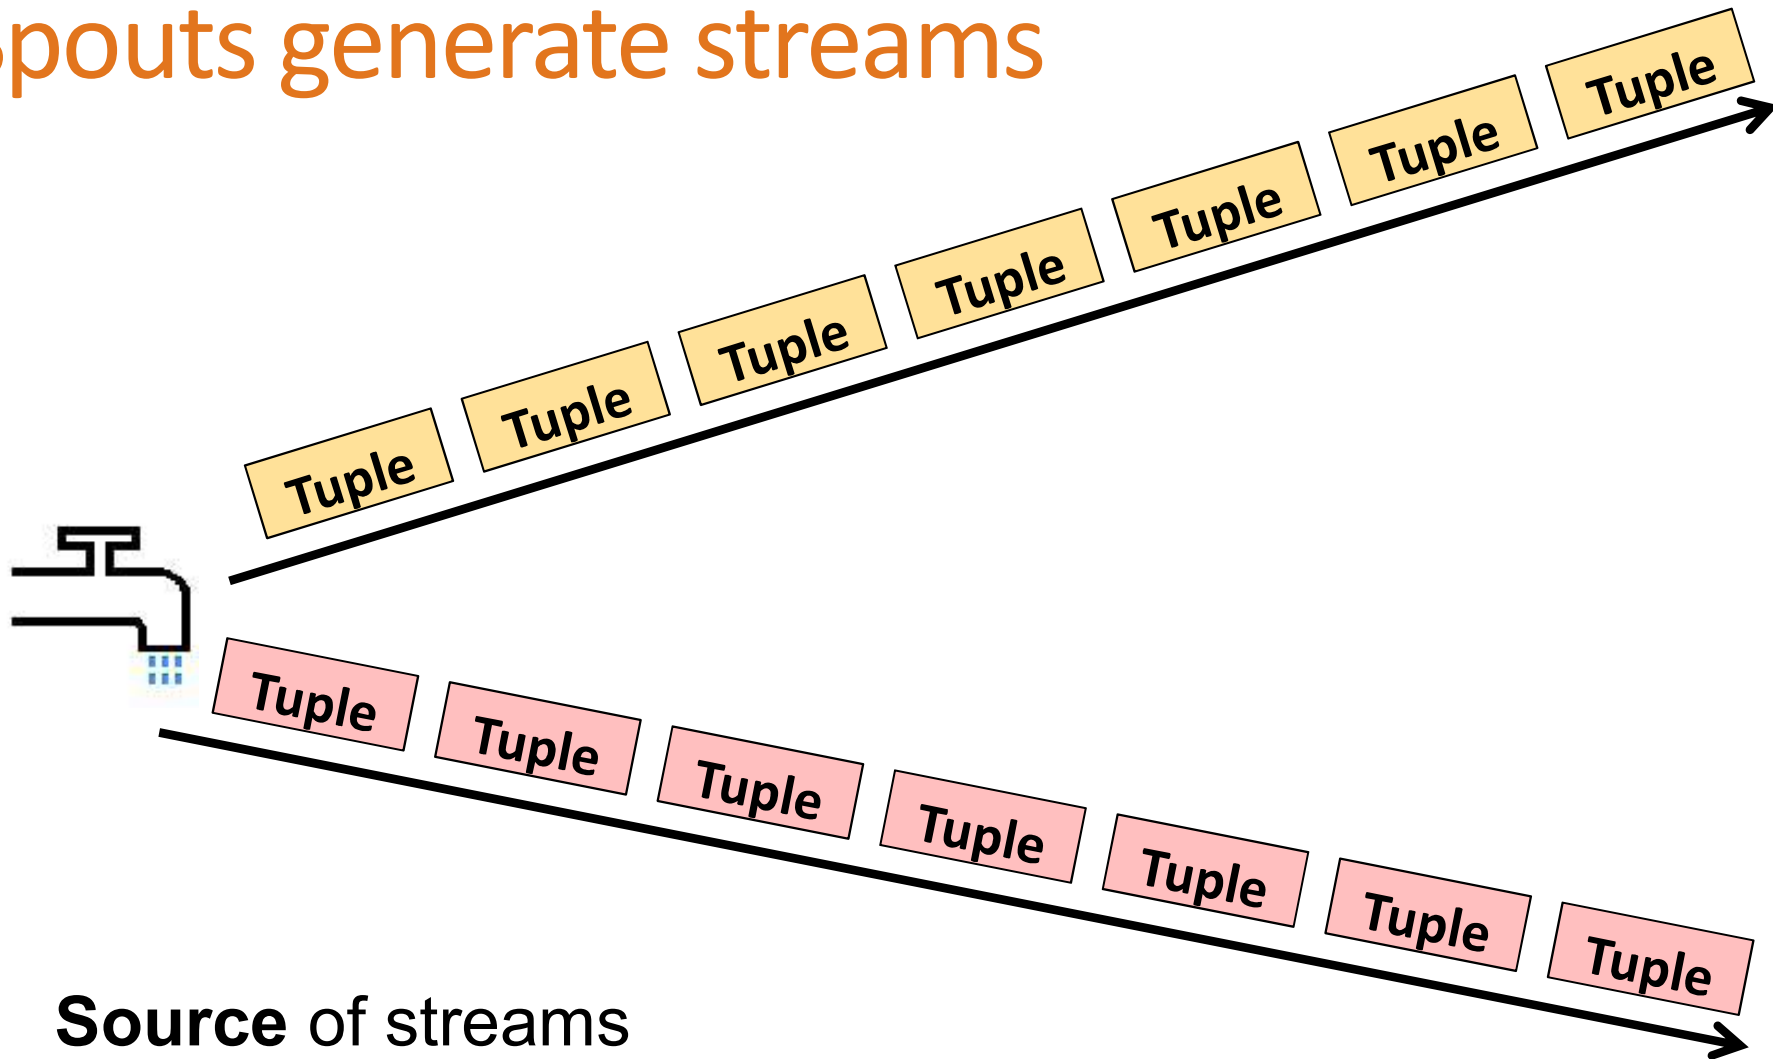
- Streams
  - Of messages arriving to bolts
- Spouts
  - Generating message streams
- Bolts
  - Consuming and generating message streams
- Topologies
  - Data flows of spouts, streams and bolts

# Streams

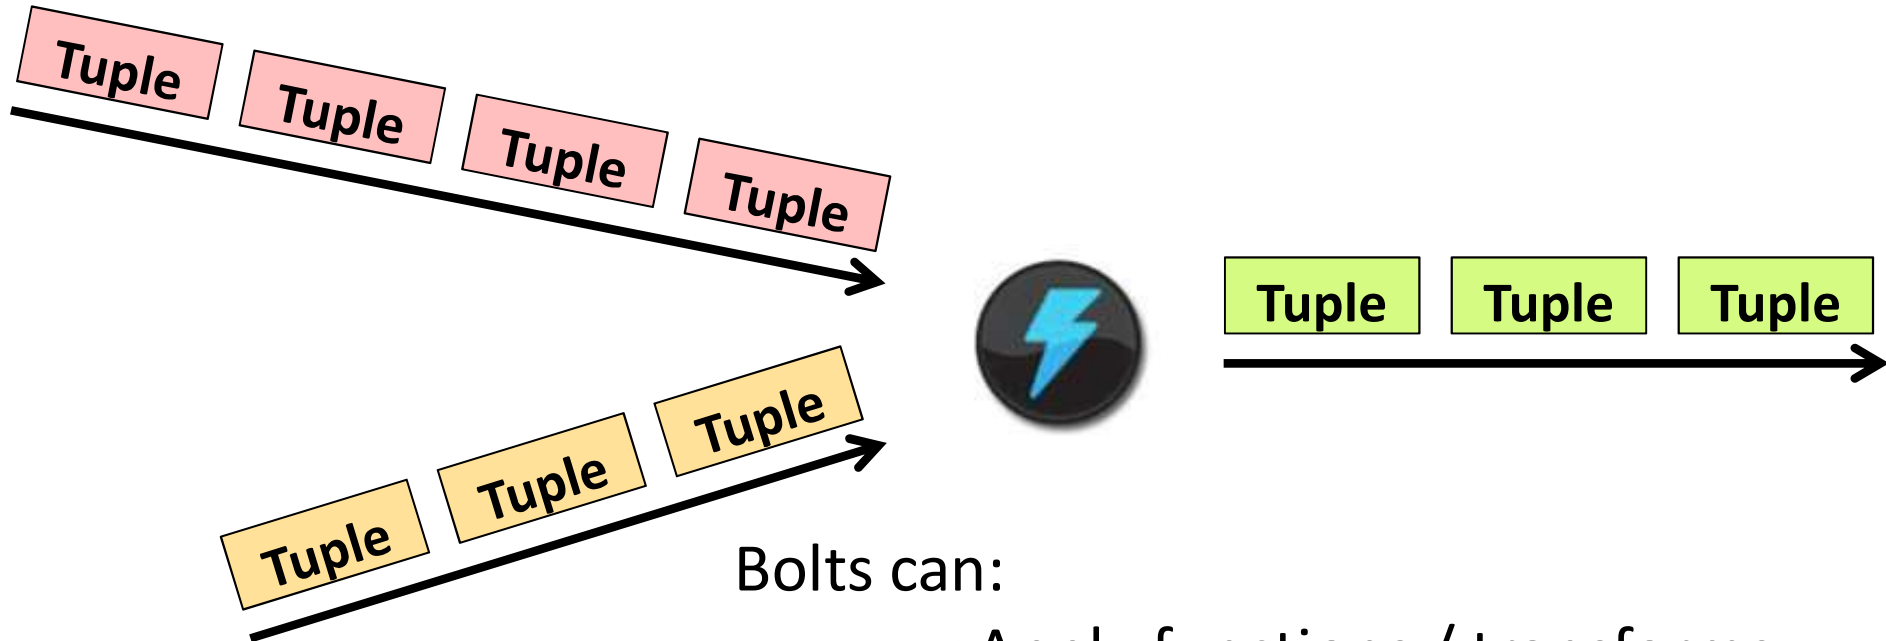


Unbounded sequence of tuples

# Spouts generate streams




# Bolts



Bolts can:

- Apply functions / transforms
- Filter
- Aggregation
- Stream joining or splitting
- Access DBs, APIs, etc

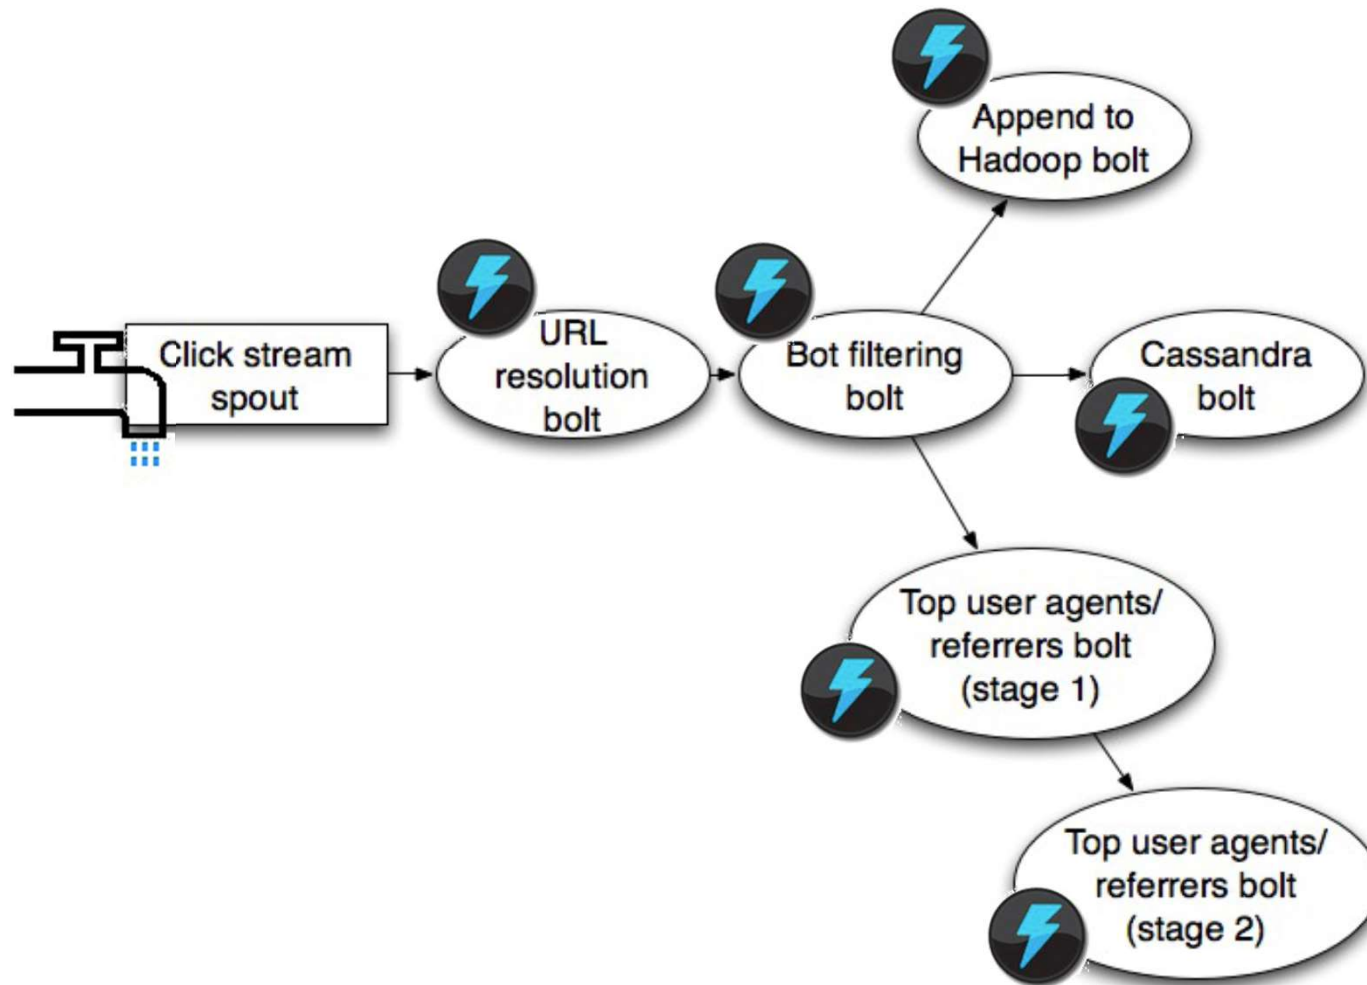
# Topology



**Network of spouts and bolts**




# Sample topology: Website click analysis



# Storm scalability

- Distribute bolts and spouts in different nodes of the cluster
- Run multiple replicas of the bolt / spout
  - Messages are grouped and sent to different replicas
    - Fields grouping: hash-based partitions
    - Random grouping: purely random, balanced
  - Problem: no longer can keep global state

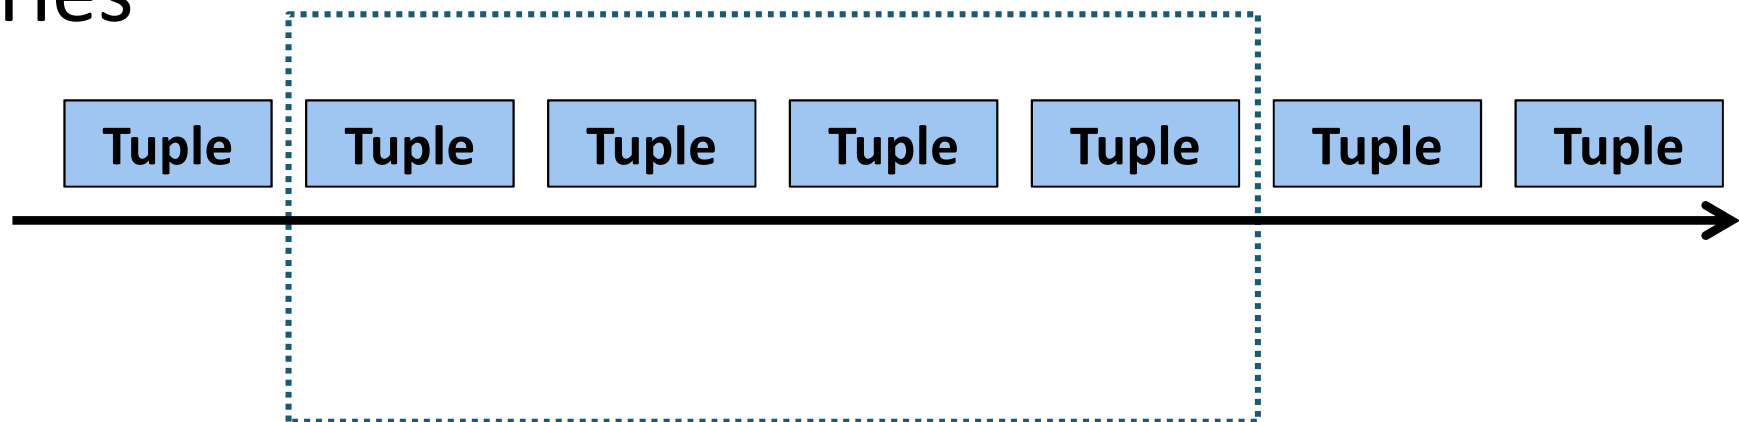
# Storm Topology




Spouts and bolts execute as  
many tasks across the cluster  
Horizontal scaling/parallelism

# Sliding Window

- Some bolts have to access more than the current tuple in order to perform its computation
- A sliding window stores a rolling list with the latest items from the stream
- Contents change over time, replaced by new entries



# Streaming word count



```
class ParseTweetBolt extends BaseBasicBolt {
```

```
  @Override
```

```
  public void execute(Tuple tuple, BasicOutputCollector collector) {
```

```
    String tweet = tuple.getString(0);
```

```
    for (String word : tweet.split(" ")) {
```

```
      collector.emit(new Values(word));
```

```
    }
```

```
  }
```

```
  @Override
```

```
  public void declareOutputFields(OutputFieldsDeclarer declarer) {
```

```
    declarer.declare(new Fields("word"));
```

```
  }
```

```
}
```

```
class WordCountBolt extends BaseBasicBolt {
```

```
  Map<String, Integer> counts = new HashMap<String, Integer>();
```

```
  @Override
```

```
  public void execute(Tuple tuple, BasicOutputCollector collector) {
```

```
    String word = tuple.getString(0);
```

```
    Integer count = counts.get(word);
```

```
    count = (count == null) ? 1 : count + 1;
```

```
    counts.put(word, count);
```

```
    collector.emit(new Values(word, count));
```

```
  }
```

```
  @Override
```


```
  public void declareOutputFields(OutputFieldsDeclarer declarer) {
```

```
    declarer.declare(new Fields("word", "count"));
```

```
  }
```

```
}
```

# Streaming word count



```
TopologyBuilder builder = new TopologyBuilder();
```

```
builder.setSpout("tweet_spout", new RandomTweetSpout(), 5);
```

```
builder.setBolt("parse_bolt", new ParseTweetBolt(), 8)
```

```
    .shuffleGrouping("tweet_spout")
```

```
    .setNumTasks(2);
```

```
builder.setBolt("count_bolt", new WordCountBolt(), 12)
```

```
    .fieldsGrouping("parse_bolt", new Fields("word"));
```

```
Config config = new Config();
```

```
config.setNumWorkers(3);
```

```
StormSubmitter.submitTopology("demo", config, builder.createTopology());
```

## tweet spout

```
class RandomTweetSpout extends BaseRichSpout {
    SpoutOutputCollector collector;
    Random rand;
    String[] tweets = new String[] {
        "@jkcrums:There's a plane in the Hudson. I'm on the ferry to pick up people. Crazy",
        "@barackobama: Four more years. pic.twitter.com/bAJE6Vom",
        ...
    };

    ....

    @Override
    public void nextTuple() {
        Utils.sleep(100);
        String tweet = tweets[rand.nextInt(tweets.length)];
        collector.emit(new Values(tweet));
    }
}
```

# Contents

- Information Streams
- Stream Processing - Storm
- **Micro batch processing - DStream**

# What is Spark Streaming?

- Framework for large scale stream processing
  - Micro-batch processing model
    - Can achieve second scale latencies
  - Scales to 100s of nodes
  - Integrates with Spark's batch and interactive processing
  - Provides a simple batch-like API for implementing complex algorithm
  - Can absorb live data streams from Kafka, Flume, ZeroMQ, etc.



# Discretized Streams

- Reuse Spark Programming model
  - Transformations on RDDs
- RDDs are created combining all the messages in a defined time interval
- A new RDD is processed at each slot
- Spark code for creating one:
  - `val streamFromMQTT = MQTTUtils.createStream(ssc, brokerUrl, topic, StorageLevel.MEMORY_ONLY_SER_2)`

# Case study: Conviva, Inc.

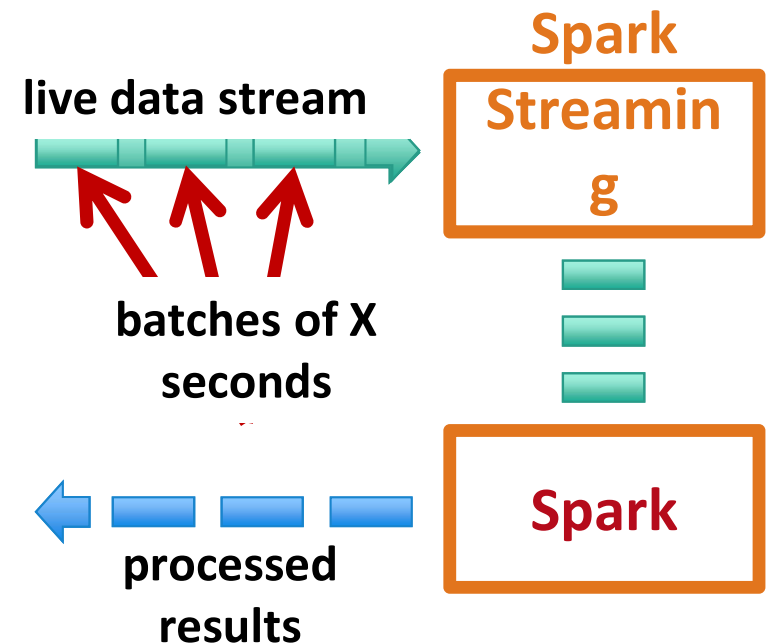
- Real-time monitoring of online video metadata
  - HBO, ESPN, ABC, SyFy, ...

- Two processing stacks
  - Custom-built distributed stream processing system
    - 1000s complex metrics on millions of video sessions
    - Requires many dozens of nodes for processing
  - Hadoop backend for offline analysis
    - Generating daily and monthly reports
    - Similar computation as the streaming system

# Discretized Stream Processing

Run a streaming computation as a **series of very small, deterministic batch jobs**

- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches



# Example 1 – Get hashtags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
```

**DStream: a sequence of RDD representing a stream of data**

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream




**stored in memory as an RDD  
(immutable, distributed)**

# Example 1 – Get hashtags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap (status => getTags(status))
```

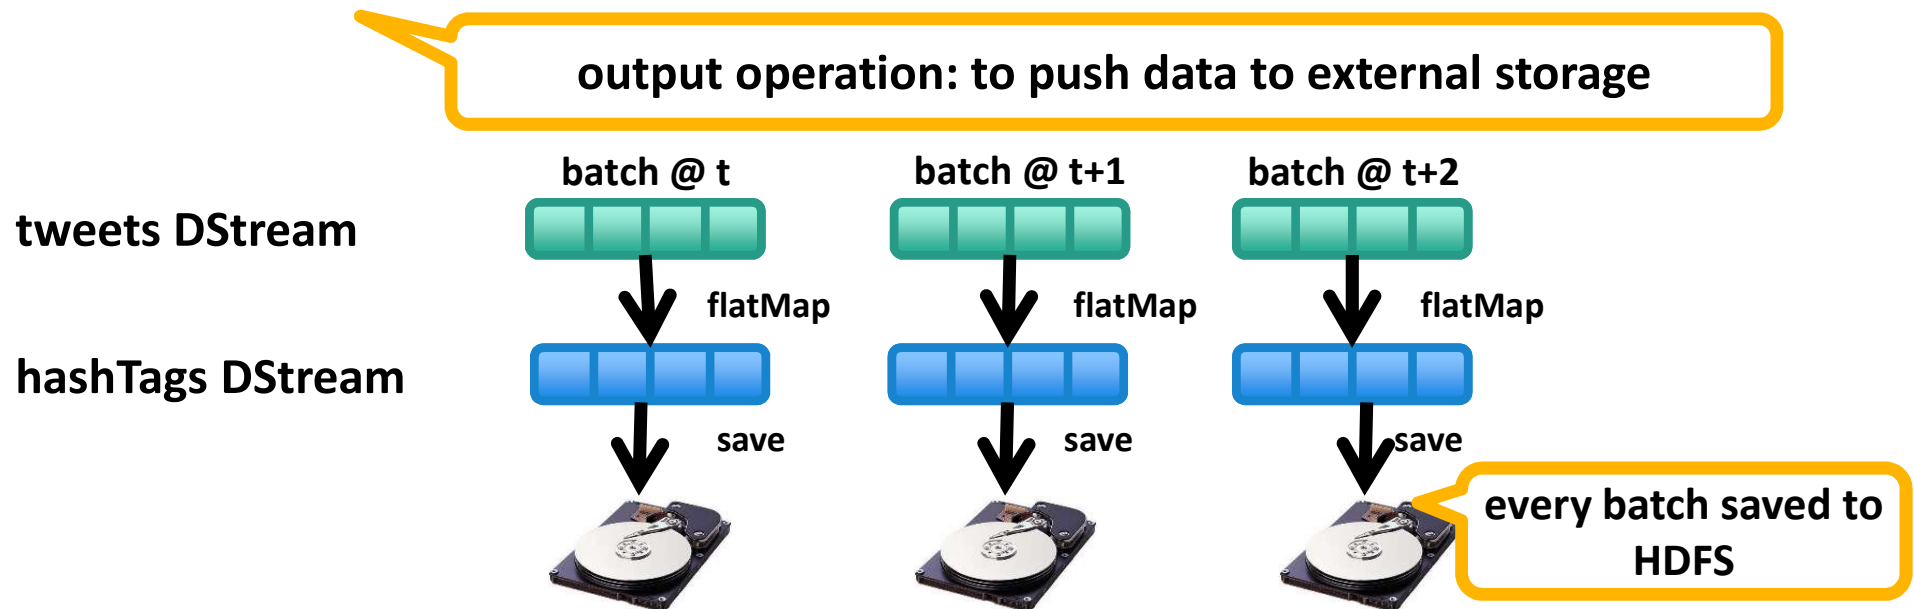
new DStream

transformation: modify data in one Dstream to create another DStream




# Example 1 – Get hashtags from Twitter

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap(status => getTags(status))
hashTags.saveAsHadoopFiles("hdfs://...")
```



## Example 2 – Count the hashtags

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)
val hashTags = tweets.flatMap(status => getTags(status))
val tagCounts = hashTags.countByValue()
```



## Example 3 – Count the hashtags over last 10 mins

```
val tweets = ssc.twitterStream(<Twitter username>, <Twitter password>)  
val hashTags = tweets.flatMap (status => getTags(status))  
val tagCounts = hashTags.window(Minutes(10), Seconds(1)).countByValue()
```



**sliding window  
operation**

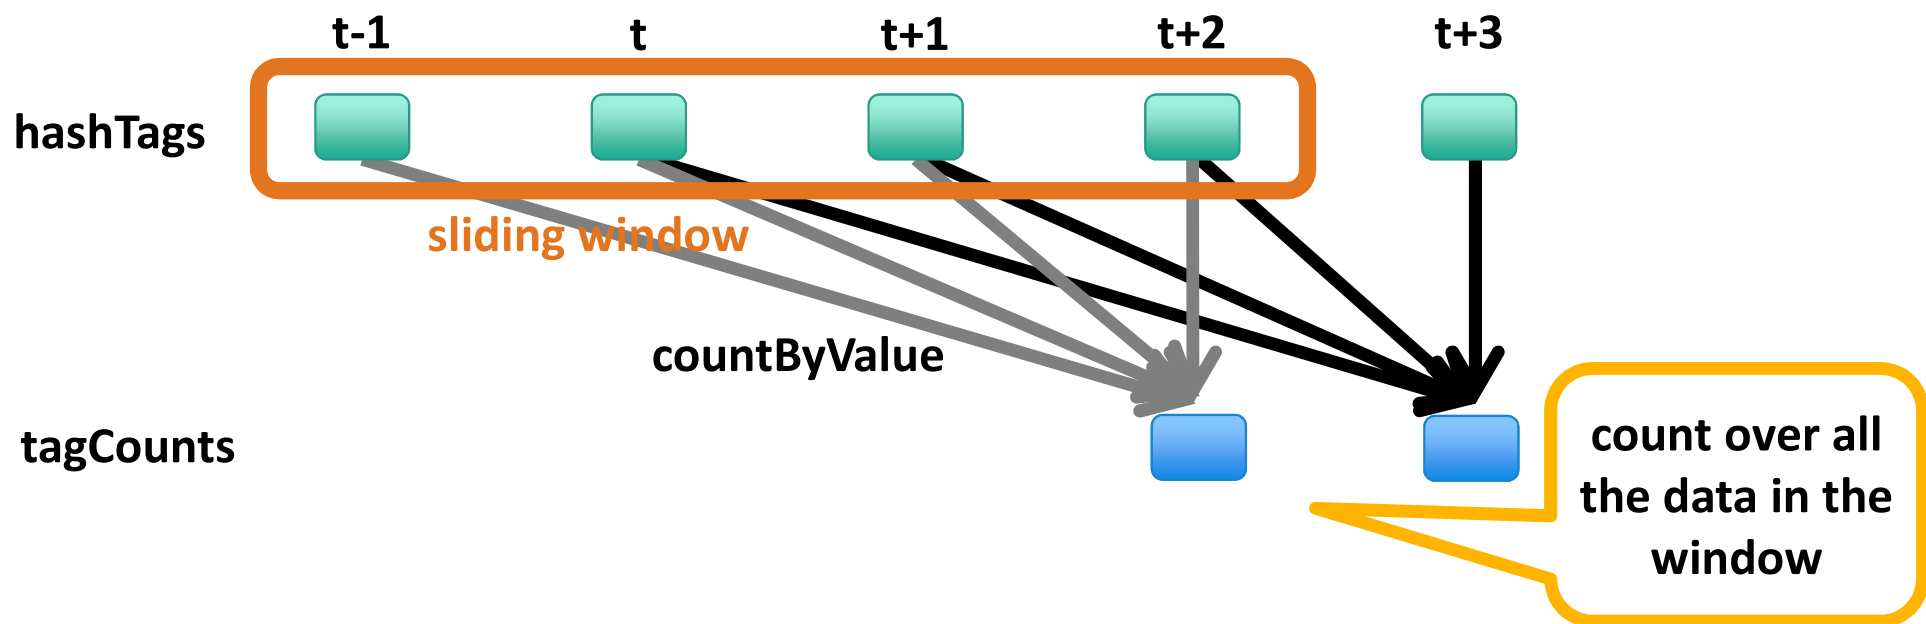
**window length**

**sliding interval**



# Example 3 – Counting the hashtags over last 10 mins

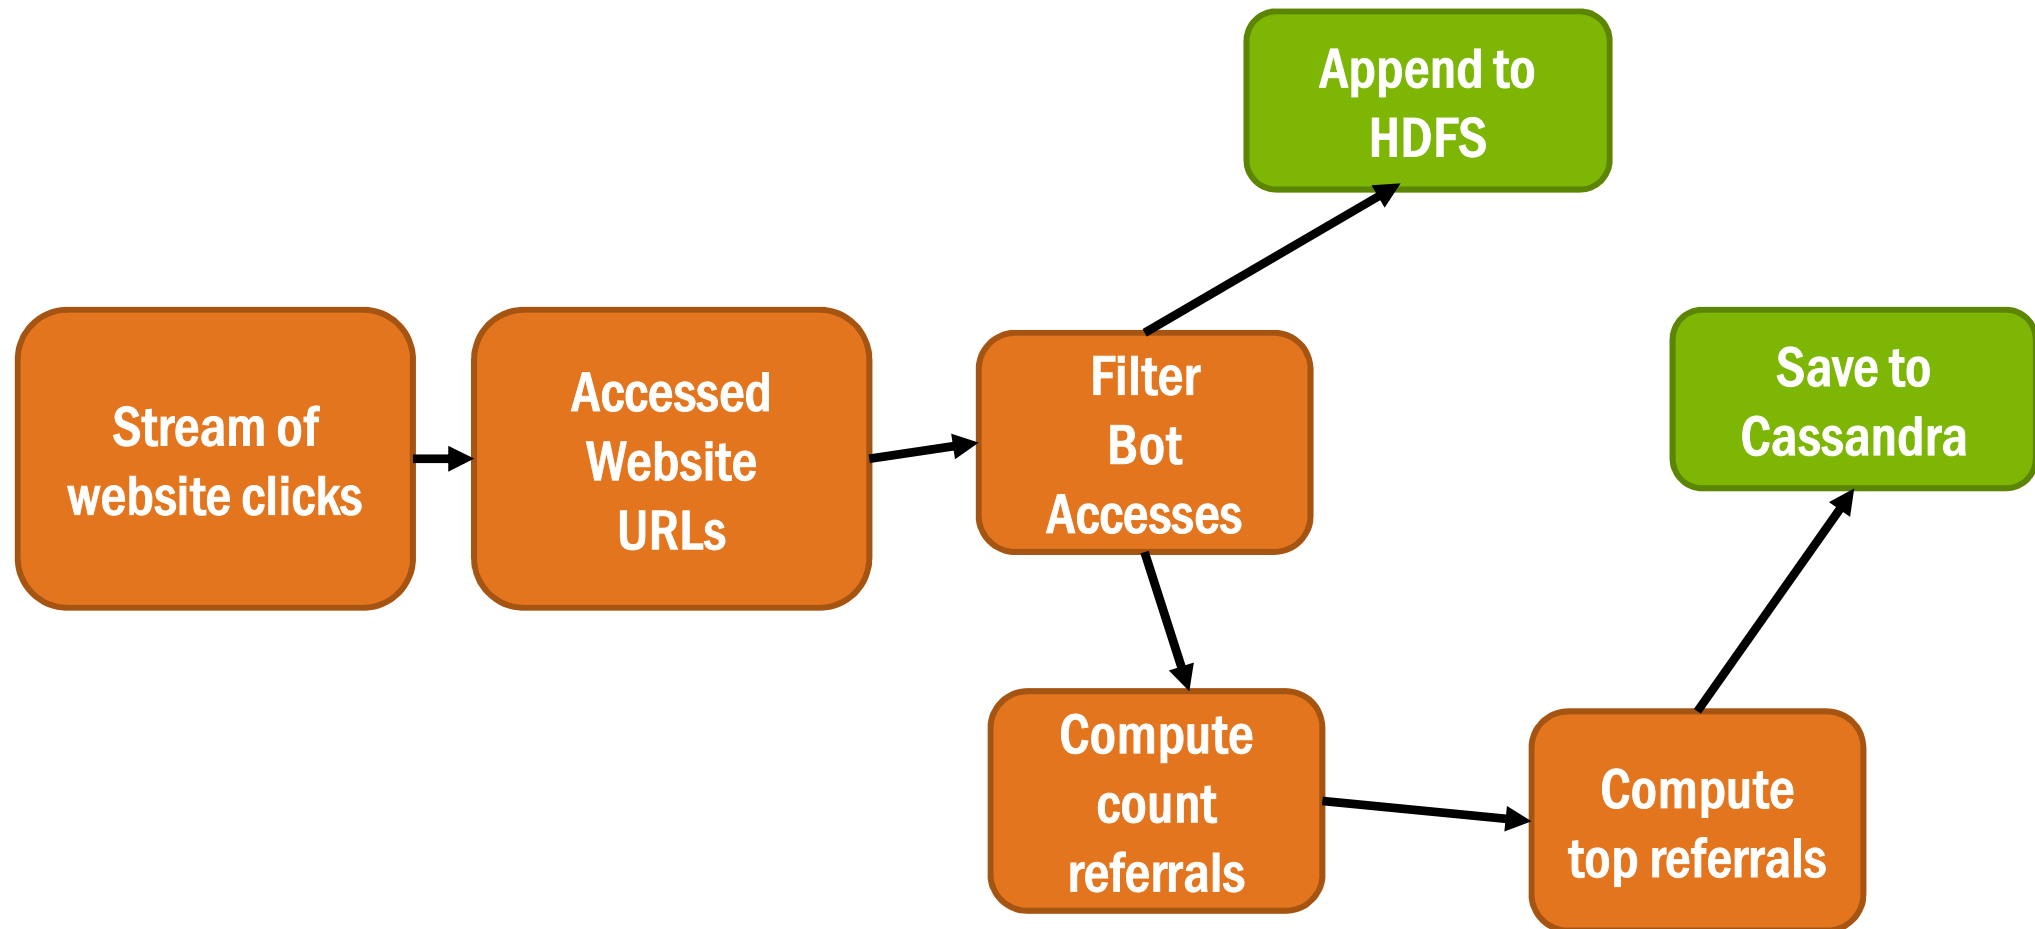
```
val tagCounts = hashTags.window(Minutes(10), Seconds(1)).countByValue()
```



## D-Stream Streaming context

- Spark streaming flows are configured by creating a `StreamingContext`, configuring what transformations flow will be done, and then invoke the `start` method
  - `val ssc = new StreamingContext(sparkUrl, "Tutorial", Seconds(1), sparkHome, Seq(jarFile))`
- There must be some action collecting in some way the results of a temporal RDD

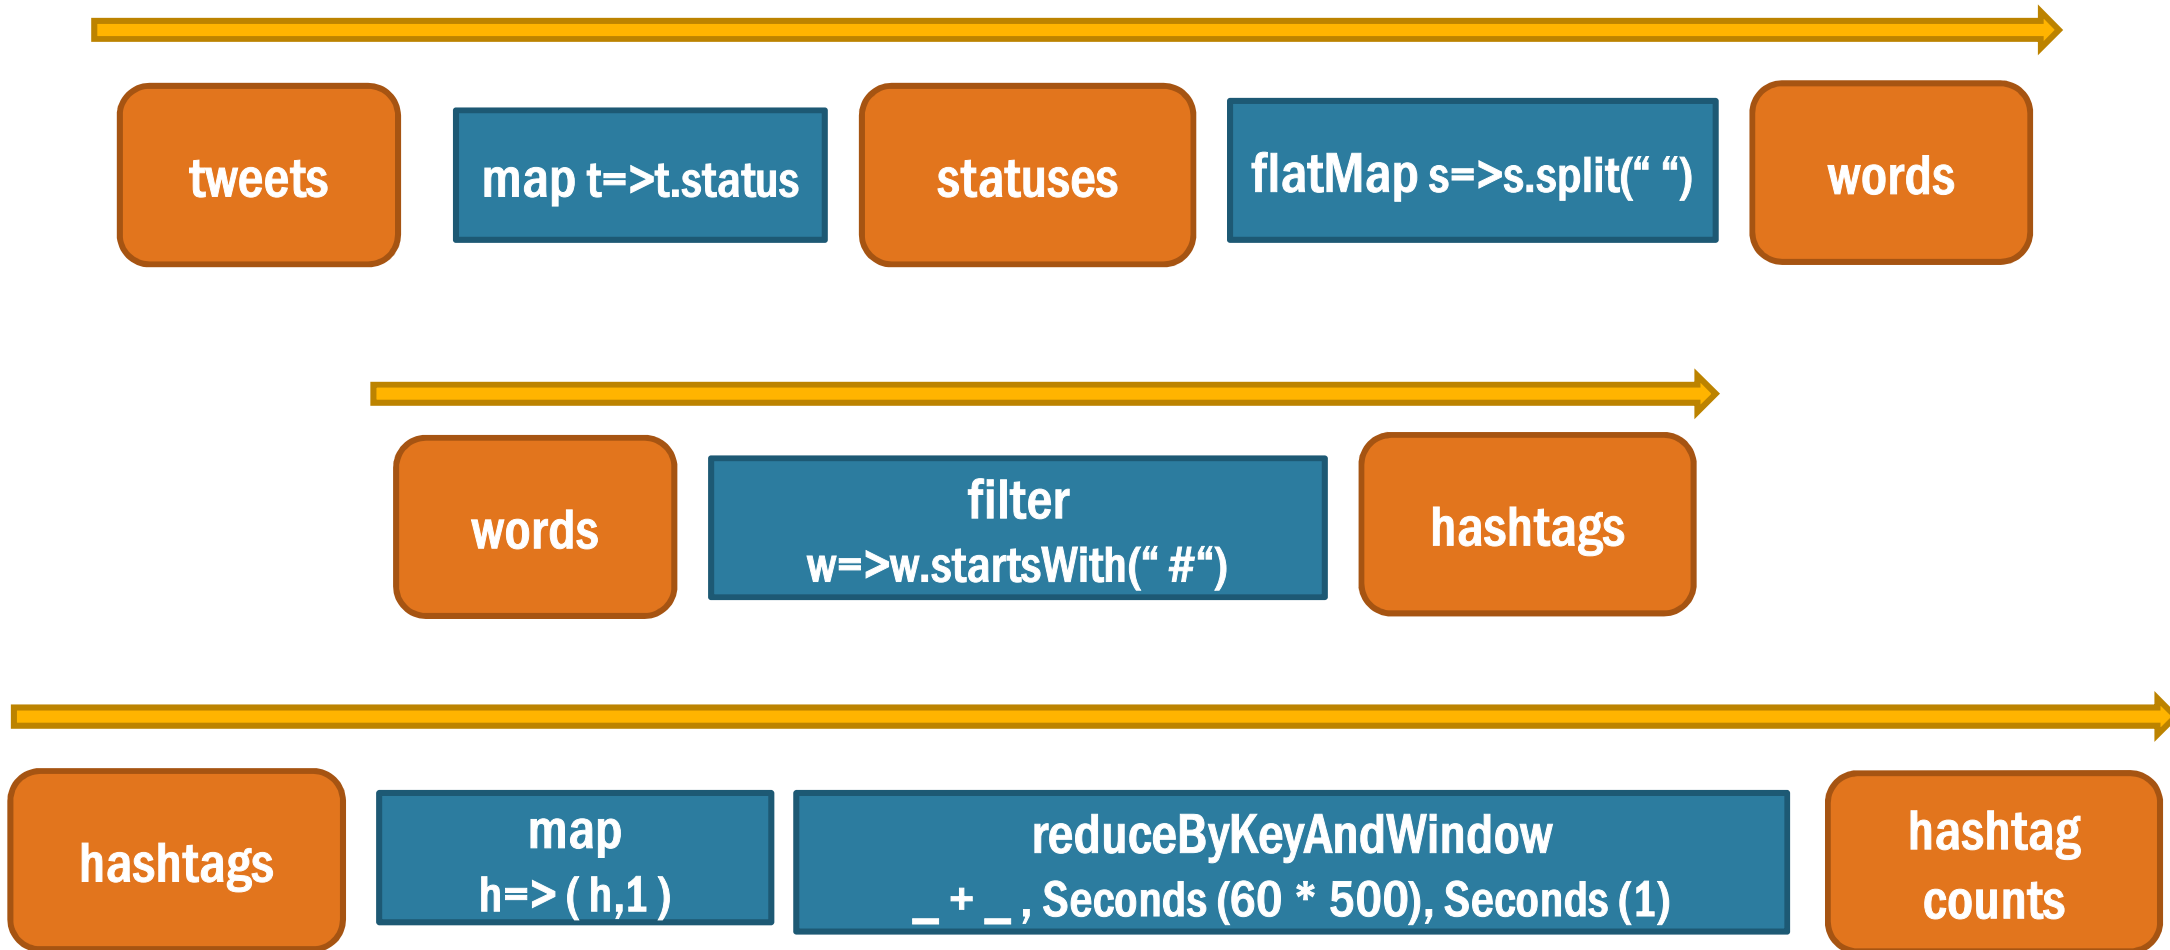
# Sample topology: Website click analysis



# Sliding window operations in Spark D-Stream

- D-Stream provides direct API support for specifying sliding windows
- Two parameters:
  - Size of the window (in seconds)
  - Frequency of computations (in seconds)
- E.g. process the maximum temperature over the last 60 seconds, every 5 seconds.
  - `reduceByWindowAndKey ( (a , b) => math . max ( a , b ) ,  
Seconds ( 60 , Seconds ( 5 ) )`

# Sample Twitter processing stream



# Sample Twitter Processing Stream

```
val ssc = new StreamingContext(sparkUrl,  
    "Tutorial", Seconds(1), sparkHome, Seq(jarFile))  
val tweets = ssc.twitterStream()  
val statuses = tweets.map(status => status.getText())  
val words = statuses.flatMap(  
    status => status.split(" "))  
val hashtags = words.filter(  
    word => word.startsWith("#"))  
val hashtagCounts = hashtags.map(tag => (tag, 1)).  
    reduceByKeyAndWindow(  
        _ + _, Seconds(60 * 5), Seconds(1))  
ssc.checkpoint(checkpointDir)  
ssc.start();
```